




# COS40003 Concurrent Programming

## Lecture 3 (b): scheduling II

# Multi-Level Feedback Queue

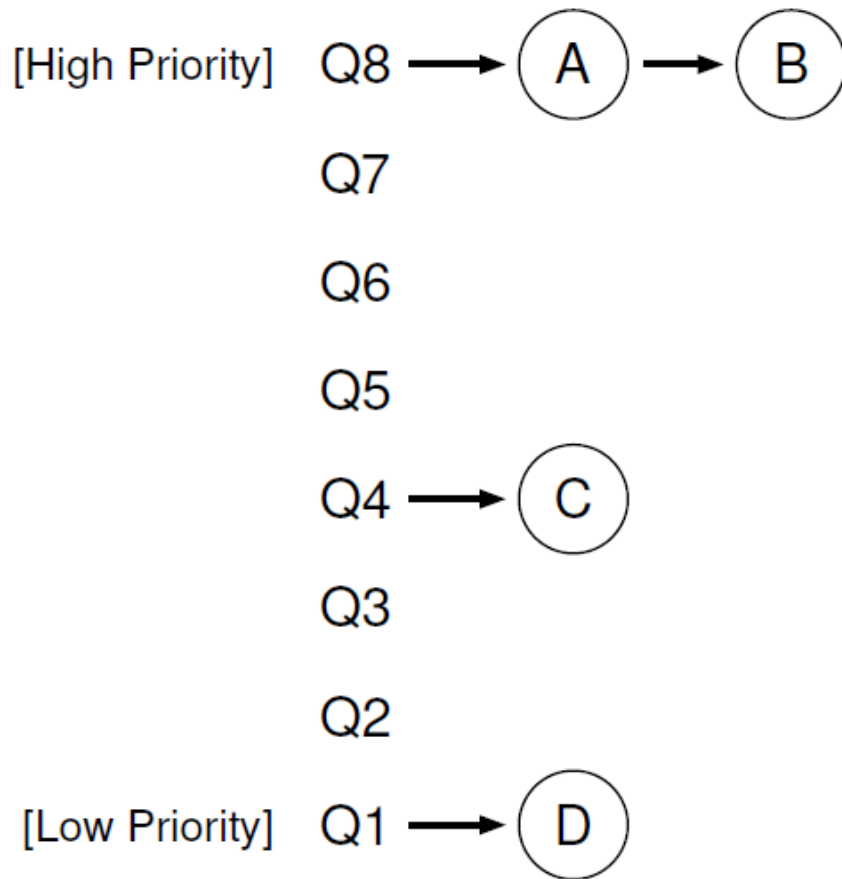
- **Fernando José Corbató**
  - **Turing Award Winner 1990**
  - "for his pioneering work in organizing the concepts and leading the development of the general-purpose, large-scale, time-sharing and resource-sharing computer systems"



# Multi-Level Feedback Queue

- Aim (achieve both)
- Optimize turnaround time
  - Dilemma: knowing SJF/PSJF is good, but how to know how long a job will run;
- Minimize response time
  - Dilemma: RR is good in response time, but is terrible for turnaround time

# MLFQ Basics




- Tasks put in **different queues**, with **different priority** level.
- **Rule 1:** If  $\text{Priority}(A) > \text{Priority}(B)$ , A runs (B doesn't).
- **Rule 2:** If  $\text{Priority}(A) = \text{Priority}(B)$ , A & B run in RR.

# MLFQ: How To Change Priority?


- Rule 3: When a job enters the system, it is placed at the highest priority (the topmost queue).
- Rule 4a: If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down to the next (lower) queue).
- Rule 4b: If a job gives up the CPU before the time slice is up, it stays at the same priority level.

# Example 1: A Single Long-Running Job



- The job enters at the highest priority (Q2) (Rule 3)
- 10ms later, the job moves to Q1 (Rule 4a)
- 10ms later, the job moves to Q0 (Rule 4a)

## Example 2: Along came a short job




- B (which will take 20ms) arrives at time  $T = 100$
- B is firstly inserted into the top queue (Rule 3)
- B completes before reaching the bottom queue, in two time slices (Rule 4a)

# How MLFQ approximates SJF?

- A job is firstly assumed to be a short job, thus giving the job high priority.
- If it actually is a short job, it will run quickly and complete;
- If it is not a short job, it will slowly move down the queues, and thus soon prove itself to be a long-running job.




# Example 3: What About I/O?




- B (an interactive job), for example, is doing a lot of I/O (say by waiting for user input from the keyboard or mouse)

# Example 3: What About I/O?



- B (an interactive job) that needs the CPU only for 1 ms before performing an I/O competing for the CPU with a long-running batch job A

# Example 3: What About I/O?



- Do not want to penalize such Task B.
- Rule 4b: if a process gives up the processor before using up its time slice, we keep it at the same priority level.

# Problems with our current MLFQ?

- It seems to do a fairly good job:
  - sharing the CPU fairly between long-running jobs
  - letting short or I/O-intensive interactive jobs run quickly. (better turnaround time)
  - Letting all jobs starting with high priority (better response time)
- **Any problem?**

# Problems with our current MLFQ?


- Starvation:
  - if there are “too many” interactive jobs in the system, they will combine to consume all CPU time, and thus long-running jobs will never receive any CPU time.
- Game the scheduler
  - before the time slice is over, issue an I/O operation (to some file you don't care about) and thus give up the CPU; doing so allows you to remain in the same queue, and thus gain a higher percentage of CPU time.

# MLFQ: How to better manage priority: Priority Boost

- To resolve starvation:
- **Rule 5:** After some time period  $S$ , move all the jobs in the system to the topmost queue.

# MLFQ: How to better manage priority: Priority Boost

- a long-running job competing for the CPU with two short-running interactive jobs




# MLFQ: How to better manage priority: better accounting

- To resolve gaming the scheduler:
- Better accounting: it does not matter CPU occupation is in one long burst or many small ones.
- **Rule 4:** Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced (i.e., it moves down one queue).  
**(To replace Rule 4(a) and Rule 4(b) )**



# Old Rules 4a and 4b VS the new anti-gaming Rule 4



# Multi-Level Feedback Queue

## Summary :

- **Rule 1:** If  $\text{Priority}(A) > \text{Priority}(B)$ , A runs (B doesn't).
- **Rule 2:** If  $\text{Priority}(A) = \text{Priority}(B)$ , A & B run in RR.
- **Rule 3:** When a job enters the system, it is placed at the highest priority (the topmost queue).
- **Rule 4:** Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced (i.e., it moves down one queue).
- **Rule 5:** After some time period  $S$ , move all the jobs in the system to the topmost queue.

# Tuning MLFQ and Other Issues

- One big question: how to **parameterize** such a MLFQ?
  - How many queues should there be?
  - How big should the time slice be per queue?
  - How often should priority be boosted?
- There are no easy answers to these questions. ☹️

# Tuning MLFQ and Other Issues

- One widely accepted rule:
  - most MLFQ variants allow for varying time-slice length across different queues
  - The high-priority queues are usually given short time slices;
  - The low-priority queues, in contrast, given longer time slices.

# Why MLfQ is good?

- Instead of demanding a priori knowledge of a job, it **observes the execution of a job and prioritizes it** accordingly
  - Has excellent overall performance (similar to SJF/PSJF) for short-running interactive jobs
  - Fair to long-running CPU-intensive workloads.
- Applied in operating systems: BSD Unix, Solaris, Windows Series

# Acknowledgement

- Chapter 7-9
  - Operating Systems: Three Easy Pieces
- 8.ppt
  - Intro to Operating System at Portland State University
  - by Jonathan Walpole

# Questions?