



# COS40003 Concurrent Programming

## Lecture 4b: Thread


# Outline

- Why we need threads?
- What is a thread?
- How threads are working?
- Thread in Java

# Java APIs

- Javadocs
  - <https://docs.oracle.com/javase/tutorial/index.html>
- Java Concurrency
  - <https://docs.oracle.com/javase/tutorial/essential/concurrency/index.html>

# Threads in Java



The `Thread` class executes instructions from its method `run()`. The actual code executed depends on the implementation provided for `run()` in a derived class.

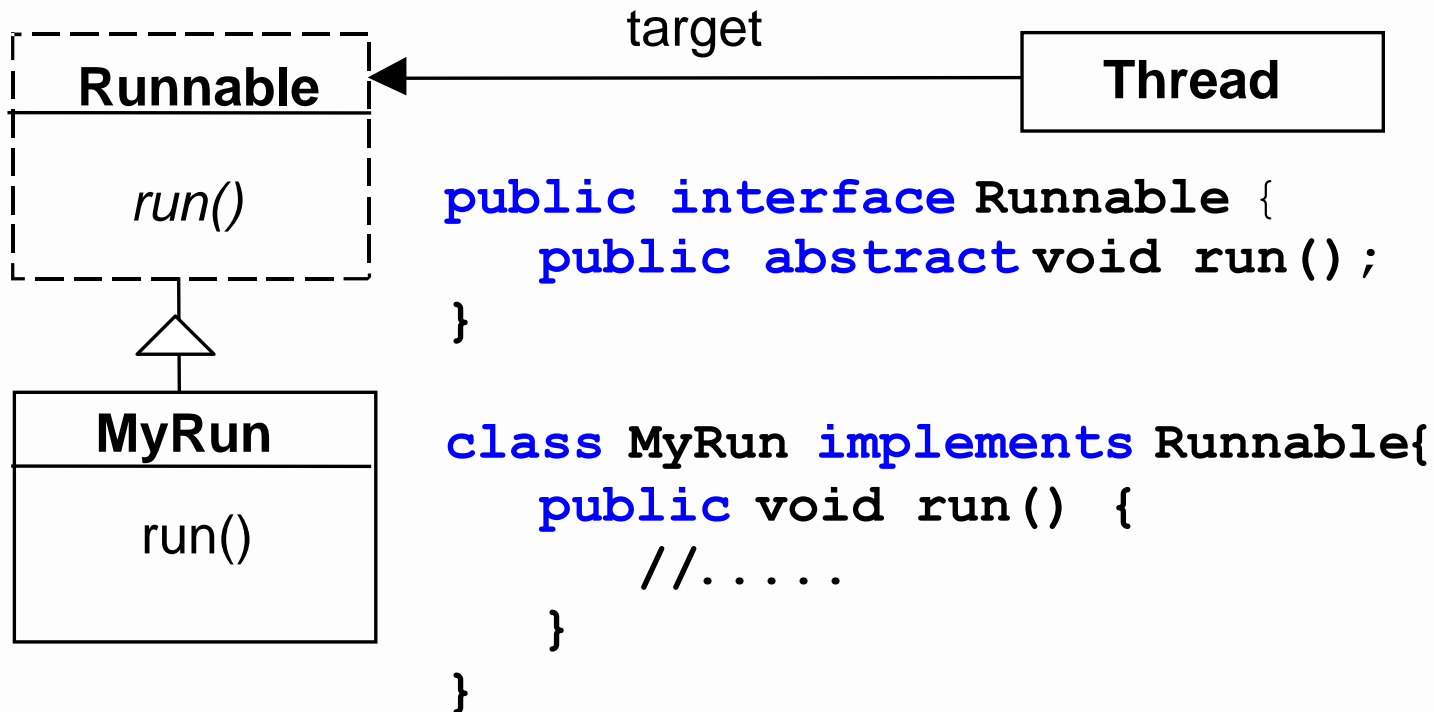
```
class MyThread extends Thread {
    public void run() {
        //.....
    }
}
```

Creating and starting a thread object:

```
Thread a = new MyThread();
a.start();
```

# Threads in Java

Since Java does not permit multiple inheritance, we often implement the **run()** method in a class not derived from Thread but from the interface Runnable. This is also more flexible and maintainable.



Creating and starting a thread object:

```
Thread b = new Thread(new MyRun());  
b.start();
```

# Threads in Java – Example 1

## The operation we want to be threaded:

```
public class PrintNumbers {  
    static public void printNumbers() {  
        for(int i=0; i<10; i++) {  
            System.out.println(  
                "Thread " +Thread.currentThread().getId() + "prints " + i);  
        }  
    }  
}
```

# Threads in Java – Example 1

## Option 1 – extending class Thread:

```
public class Thread1 extends Thread {  
  
    @Override  
    public void run() {  
        System.out.println("Thread Id: " +  
Thread.currentThread().getId());  
        // do our thing  
        PrintNumbers.printNumbers();  
    }  
}
```

# Threads in Java – Example 1

## Option 1 – extending class Thread (cont’):

```
public static void main(String[] args) {  
    System.out.println("Main ThreadId: " +  
        Thread.currentThread().getId());  
    for(int i=0; i<3; i++) {  
        new Thread1().start(); // don't call run!  
    }  
    PrintNumbers.printNumbers();  
}
```



# Threads in Java – Example 1

## Option 2 – implementing Runnable:

```
public class Thread2 implements Runnable {  
  
    @Override  
    public void run() {  
        System.out.println("Thread Id: " +  
            Thread.currentThread().getId());  
        // do our thing  
        PrintNumbers.printNumbers();  
    }  
}
```

# Threads in Java – Example 1

## Option 2 – implementing Runnable (cont’):

```
public static void main(String[] args) {  
    System.out.println("Main ThreadId: " +  
        Thread.currentThread().getId());  
    for(int i=0; i<3; i++) {  
        new Thread(new Thread2()).start();  
        // again, don't call run!  
    }  
    PrintNumbers.printNumbers();  
}
```

# Creating Threads

- extending the Thread class
  - must implement the *run()* method
  - thread ends when *run()* method finishes
  - call *.start()* to get the thread ready to run

# Advantage of Using Runnable

- remember - can only extend one class in Java
- implementing runnable allows class to extend something else

# Some useful functions

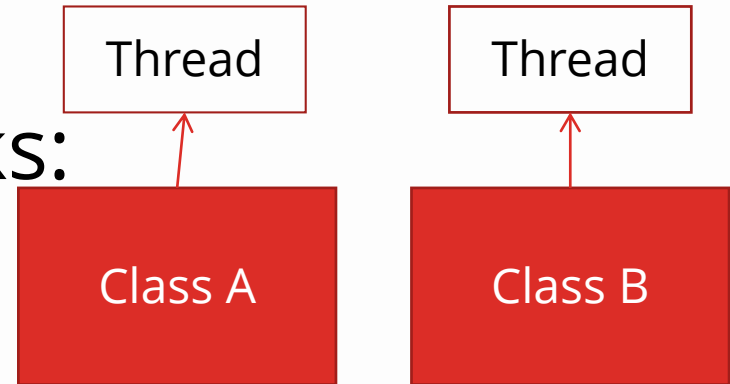
- *\_.start()*: begins a thread running
- *wait()* and *notify()*: for synchronization
  - more on this later
- *\_.stop()*: kills a specific thread (deprecated)
- *\_.suspend()* and *resume()*: deprecated
- *\_.join()*: wait for specific thread to finish
- *\_.setPriority()*: 0 to 10 (MIN\_PRIORITY to MAX\_PRIORITY); 5 is default (NORM\_PRIORITY)

## Example 2 – concurrent independent tasks

- Sometimes we just want to run independent actions concurrently and (possibly) benefit from multi-cores

## Example 2 – concurrent independent tasks

- Two concurrent tasks:



```
public void myFunction() {  
    Thread task1 = new A();  
    Thread task2 = new B();  
  
    task1.start();  
    task2.start();  
  
    task1.join();  
    task2.join();  
}
```

# Example 2 – concurrent independent tasks

```
public class A extends Thread{
    public void run()
    {
        while(true)
        {
            System.out.println("This is Thread A");
            try{
                Thread.sleep(1000);
            }catch(InterruptedException e){}
        }
    }
}
```

```
public static void main() {

    Thread t1 = new A();
    Thread t2 = new B();

    t1.start();
    t2.start();

}
```

```
public class B extends Thread{
    public void run()
    {
        while(true)
        {
            System.out.println("This is Thread B");
            try{
                Thread.sleep(3000);
            }catch(InterruptedException e){}
        }
    }
}
```



# Thread safety in Java libraries

- Many of the Java library classes are *not* thread safe!
- Classes in **java.lang** are thread safe
- Classes in **java.util** might **not** be safe! Check the documentation before use.
- Examples:
  - `ArrayList` and other collections from `java.util` are not thread safe; two threads changing the same list at once may break it.
  - Java GUIs are not thread safe; if two threads are modifying a GUI simultaneously, they may put the GUI into an invalid state.
- Counterexamples:
  - The `Random` class chooses numbers in a thread-safe way.
  - Some input/output (like `System.out`) is thread safe.

# Example 3 – Shared variables in concurrency

- Atomicity
  - An action is atomic if it is indivisible
  - In Java, integer increment is not atomic

```
i++;
```

is actually

1. Load data from variable *i*
2. Increment data by 1
3. Store data to variable *i*

# Example 3 (cont.)

```
3 public class IncrementTest implements Runnable{
4
5     static int classData = 0;
6     int instanceData = 0;
7
8     @Override
9     public void run() {
10         int localData = 0;
11
12         while (localData < 100000000) {
13             localData++;
14             instanceData++;
15             classData++;
16         }
17
18         System.out.println("localData: " + localData +
19                             "\tinstanceData: " + instanceData +
20                             "\tclassData: " + classData);
21     }
22
23
24     public static void main(String[] args) {
25         // TODO Auto-generated method stub
26         IncrementTest instance = new IncrementTest();
27
28         Thread t1 = new Thread(instance);
29         Thread t2 = new Thread(instance);
30
31         t1.start();
32         t2.start();
33     }
34 }
35
36 19
37 }
```

shared between all instances of this class

shared between functions of an object

not shared with anyone


## Task

- What is the output of this program with a single thread? Why?
- How about with two threads? Why?

# Results

- localData = 10,000,000
- instanceData != 20,000,000
- classData != 20,000,000
- Why?
- Try: create thread t2 with another instance and see the result:

```
IncrementTest instance2 = new IncrementTest();  
Thread t2 = new Thread(instance2);
```




# Thread Pools

- A program that creates a **huge number of short-lived** threads can be **inefficient**.
  - Threads are managed by OS. There is a **cost** for creating them.
  - Cost/overhead = memory + time
- The cost can be reduced by Thread Pools
  - A Thread Pool creates a number of threads and keeps them alive. (It asks for memory all at once not on demand).

# Client-Server Example

Why it is not good to create a thread upon each request?

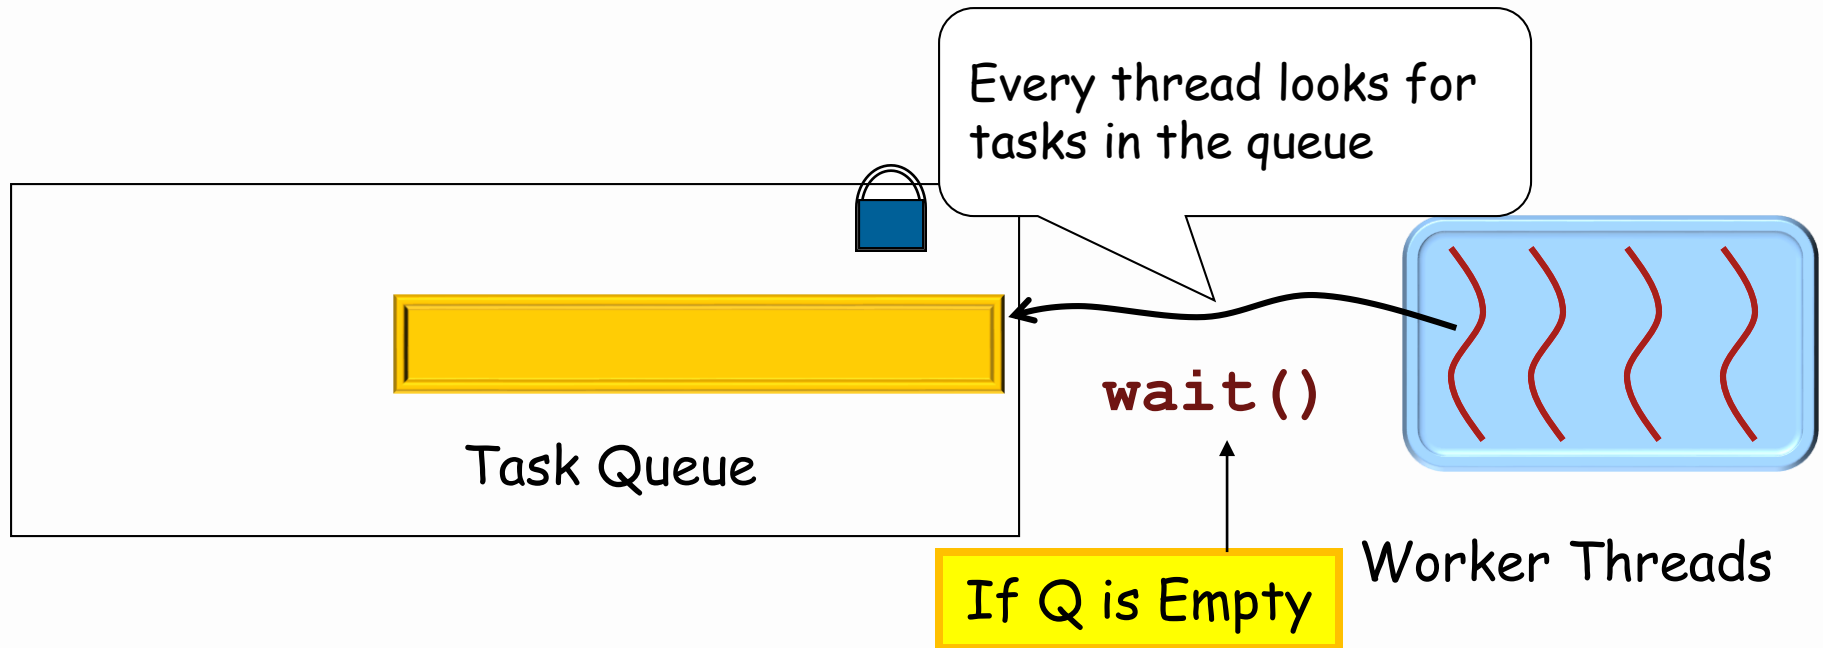


# Thread Pool, sever Example

- Thread pools are especially important in client-server applications
  - The processing of each individual task is short-lived and the number of requests is large
  - Servers should not spend more time and consume more system resources creating and destroying threads than processing actual user requests
- **When too many requests arrive, thread pools enable the server to force clients to wait until threads are available**





# How Thread Pools work



All the worker threads wait for tasks

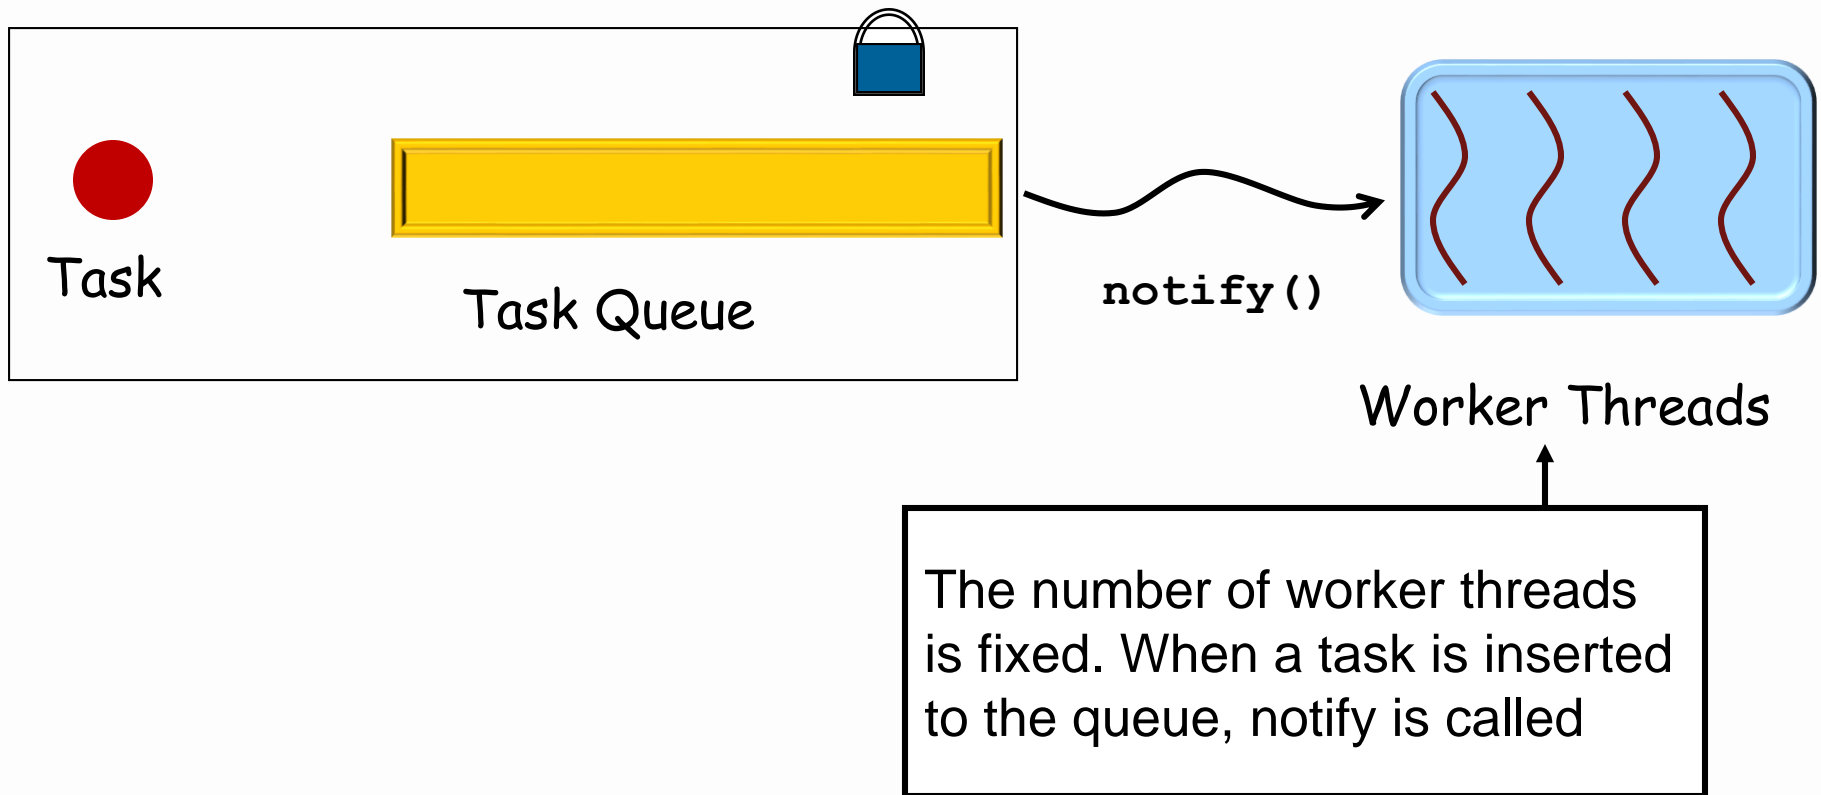
# How Thread Pools work



Worker Threads



The number of worker threads is fixed. When a task is inserted to the queue, notify is called

# How Thread Pools work



# How Thread Pools work



The task is executed by  
the thread



Worker Threads

# How Thread Pools work

The task is executed by  
the thread





Worker Threads


The remaining tasks are  
executed by the other threads

# How Thread Pools work

When a task ends, the thread is released



Task Queue





Worker Threads

While the Q is not empty, take the task from the Q and run it (if the Q was empty, wait() would have been called)

# How Thread Pools work

A new task is executed  
by the released thread



Worker Threads

# Thread Pools in Java, example 1

(Task class implements Runnable)

```
Task task1 = new Task();  
Task task2 = new Task();  
Task task3 = new Task();
```

```
System.out.println("Starting threads");
```

```
ExecutorService executor = Executors.newFixedThreadPool(3);
```

```
executor.execute(task1);  
executor.execute(task2);  
executor.execute(task3);  
executor.shutdown(); // shutdown worker threads  
executor.awaitTermination(1, TimeUnit.NANOSECONDS);
```



# What's happening

- The code in method `main` executes in the main thread
- Creates three Task objects (no extra threads are in this stage)
- Create an executor that uses a fixed thread pool by invoking **`newFixedThreadPool()`** method in `Executors` (returns a `ExecutorService` object)
- **`execute()`** creates a new Thread inside the *ExecutorService* and returns immediately from each invocation
- *ExecutorService* method **`shutdown()`** Initiates an orderly shutdown in which previously submitted tasks are executed, but no new tasks will be accepted.
- **`awaitTermination()`** waits for submitted tasks to finish
- The program will not terminate until its last thread completes execution

# Java Thread Pool APIs

- Thread Pool Tutorial
  - <https://docs.oracle.com/javase/tutorial/essential/concurrency/pools.html>
- `java.util.concurrent.ThreadPoolExecutor`
  - <https://docs.oracle.com/javase/8/docs/api/java/util/concurrent/ThreadPoolExecutor.html>
  - `execute()`, `shutdown()`, `awaitTermination()` are within

# Pool Size

- What is better: to have a large pool or a small pool?
- Each thread consumes resources
  - memory, management overhead, etc.
  - A large pool can cause *starvation*
- Incoming tasks wait for a free thread
  - A small pool can cause *starvation*
- Therefore, you have to *tune the thread pool size* according to the number and characterizations of expected tasks
- There should also be a limit on the size of the task queue (*why?*)

# Handling too Many Requests

- **What is the problem with the server being overwhelmed with requests?**
- What can a server do to avoid a request overload?
  - **Do not add to the queue** all the requests: ignore or send an error response
  - **Use several pool sizes** according to stress characteristics (but do not change the size too often...)

# Acknowledgement

- Chapter 26
  - Operating Systems: Three Easy Pieces
- 3.ppt
  - Intro to Operating System at Portland State University
  - by Jonathan Walpole

# Questions?