



# COS40003 Concurrent Programming


## Lecture 5: Lock (I)

# Review

- Question: why process context switch is more expensive than thread context switch?
- Process switch:
  - CPU state and memory working space
- Thread switch
  - CPU state

# Recall: process context switch

- Save registers, program counter, stack pointer of Prog1



# Outline

- Why there is concurrency problem?
- Key concurrency terms
- Lock implementation
  - Using Lock object
  - Using Synchronized keyword

# Example 3 (Recall)

```
3 public class IncrementTest implements Runnable{
4
5     static int classData = 0;
6     int instanceData = 0;
7
8     @Override
9     public void run() {
10         int localData = 0;
11
12         while (localData < 100000000) {
13             localData++;
14             instanceData++;
15             classData++;
16         }
17
18         System.out.println("localData: " + localData +
19                             "\tinstanceData: " + instanceData +
20                             "\tclassData: " + classData);
21     }
22
23
24     public static void main(String[] args) {
25         // TODO Auto-generated method stub
26         IncrementTest instance = new IncrementTest();
27
28         Thread t1 = new Thread(instance);
29         Thread t2 = new Thread(instance);
30
31         t1.start();
32         t2.start();
33     }
34 }
35
36
37
```

is shared between all instances of this class

If done in a single procedure

```
{
IncrementTest t1 =
new IncrementTest()
```

```
IncrementTest t2 =
new IncrementTest()
```

```
t1.classData ++;
t2.classData ++;
```

```
// classData increased
twice
}
```

# Example 3 (Recall)

```
3 public class IncrementTest implements Runnable{
4
5     static int classData = 0;
6     int instanceData = 0;
7
8     @Override
9     public void run() {
10         int localData = 0;
11
12         while (localData < 100000000) {
13             localData++;
14             instanceData++;
15             classData++;
16         }
17
18         System.out.println("localData: " + localData +
19                             "\tinstanceData: " + instanceData +
20                             "\tclassData: " + classData);
21     }
22
23
24     public static void main(String[] args) {
25         // TODO Auto-generated method stub
26         IncrementTest instance = new IncrementTest();
27
28         Thread t1 = new Thread(instance);
29         Thread t2 = new Thread(instance);
30
31         t1.start();
32         t2.start();
33     }
34 }
35
36
37
```

Each object has a copy of instanceData

```
{
IncrementTest t1 =
new IncrementTest()

IncrementTest t2 =
new IncrementTest()

t1.instanceData ++;
t2.instanceData ++;

// t1's instanceData
increased once
// t2's instanceData
increased once
}
```

# Example 3 (Recall)

```
3 public class IncrementTest implements Runnable{
4
5     static int classData = 0;
6     int      instanceData = 0;
7
8     @Override
9     public void run() {
10         int localData = 0;
11
12         while (localData < 100000000) {
13             localData++;
14             instanceData++;
15             classData++;
16         }
17
18         System.out.println("localData: " + localData +
19                             "\tinstanceData: " + instanceData +
20                             "\tclassData: " + classData);
21     }
22
23
24     public static void main(String[] args) {
25         // TODO Auto-generated method stub
26         IncrementTest instance = new IncrementTest();
27
28         Thread t1 = new Thread(instance);
29         Thread t2 = new Thread(instance);
30
31         t1.start();
32         t2.start();
33
34     }
35
36 }
37
```

not shared, used within the calling function

```
{
    IncrementTest t1 =
    new IncrementTest()

    IncrementTest t2 =
    new IncrementTest()

    t1.start();
    t1.start();

    // localData only
    available in the function

}
```

# Results

- localData = 10,000,000
- instanceData != 20,000,000
- classData != 20,000,000
- Why?
- Try: create thread t2 with another instance and see the result:

```
IncrementTest instance2 = new IncrementTest();  
Thread t2 = new Thread(instance2);
```



# Example 3 – Shared variables in concurrency

`i++;` is actually

1. Load data from variable `i`
2. Increment data by 1
3. Store data to variable `i`

- Decomposed in assembly code:
  1. Mov Addr1 R1 (mov value at Addr1 to R1)
  2. Add 1 R1 (increase Register R1 by 1)
  3. Mov R1 Addr1 (store value at R1 to Addr1)

Addr1: value in memory

R1: value in CPU register

# Beginning value at Addr1 is 50

• Thread1	• Thread2	R1	Addr1
1. Mov Addr1 R1		50	50
2. Add 1 R1		51	50
(context switch to Thread 2, R1 saved)	3. Mov Addr1 R1	50	50
	4. Add 1 R1	51	50
(context switch back to Thread1, R1 restored)	5. Mov R1 Addr1	51	51
6. Mov R1 Addr1		51	51



# Outline

- Why there is concurrency problem?
- Key concurrency terms
- Lock implementation
  - Using Lock object
  - Using Synchronized keyword

# Key Concurrency Terms

- Critical section
- Race condition
- Indeterminate
- Mutual exclusion
- Atomicity

# Critical section


- A **critical section** is a piece of code that accesses a shared resource, usually a variable or data structure.

```
public void run() {  
    int localData = 0;  
  
    while (localData < 10000000) {  
        localData++;
```

```
        instanceData++;  
        classData++;
```

```
    }
```





Attribute to: The lord of the rings

# Race condition

- A **race condition** arises if multiple threads of execution enter the critical section at roughly the same time; eg., both attempt to update the shared data structure, leading to an undesirable outcome.
- Example:
  - Both threads want to run  
instanceData ++  
classData++

# Indeterminate

- An **indeterminate** program consists of one or more race conditions; the output of the program varies from run to run, depending on which threads ran when.
- Example:
  - Run example3 multiple times, get different results



# Mutual exclusion

- To avoid these problems, threads should use some kind of **mutual exclusion** primitives; doing so guarantees that only a single thread ever enters a critical section, thus avoiding races, and resulting in deterministic program outputs.
- Example
  - Lock (introduced later)

# Atomicity

- Means: all or none  
either it has not run at all, or it has run to completion

Similar idea applied in database:

- Transaction: grouping of many actions into a single atomic action
  - Eg., A transfers \$50 to B


# Edsger W. Dijkstra



- Turing Award (1972)
- “No other individual has had a larger influence on research in principles of distributed computing”
- [https://en.wikipedia.org/wiki/Edsger\\_W.\\_Dijkstra](https://en.wikipedia.org/wiki/Edsger_W._Dijkstra)

# Dijkstra's algorithm

- is an algorithm for finding the shortest paths between nodes in a graph




# Key Concurrency Terms Summary

- A **critical section** is a piece of code that accesses a *shared* resource, usually a variable or data structure.
- A **race condition** arises if multiple threads of execution enter the critical section at roughly the same time; eg., both attempt to update the shared data structure, leading to an undesirable outcome.
- An **indeterminate** program consists of one or more race conditions; the output of the program varies from run to run, depending on which threads ran when.
- To avoid these problems, threads should use some kind of **mutual exclusion** primitives; doing so guarantees that only a single thread ever enters a critical section, thus avoiding races, and resulting in deterministic program outputs.


# Outline

- Why there is concurrency problem?
- Key concurrency terms
- Lock implementation
  - Using Lock object
  - Using Synchronized keyword

# How to guarantee ME (Mutual Exclusion)?



## Critical Section



# A Lock is similar to a phone booth

- Phone booth can accommodate only one person
- If the booth is empty the first person goes inside
- 2nd person has to wait until the first person leaves the booth



Bangkok ,1986



# Acquiring and Releasing Locks

**Thread B**

**Thread C**

**Thread A**

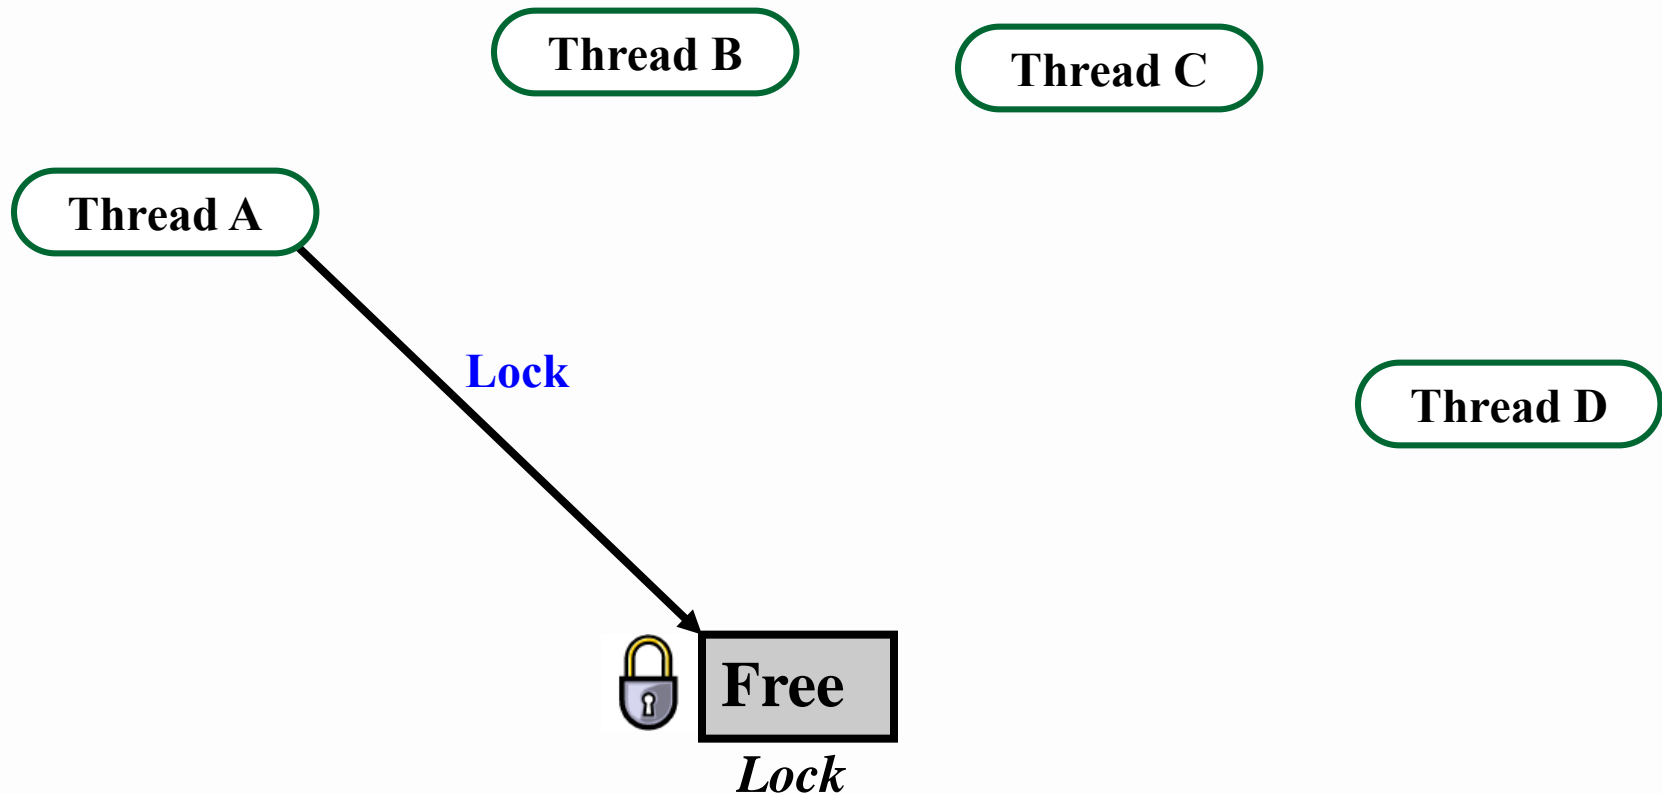
**Thread D**



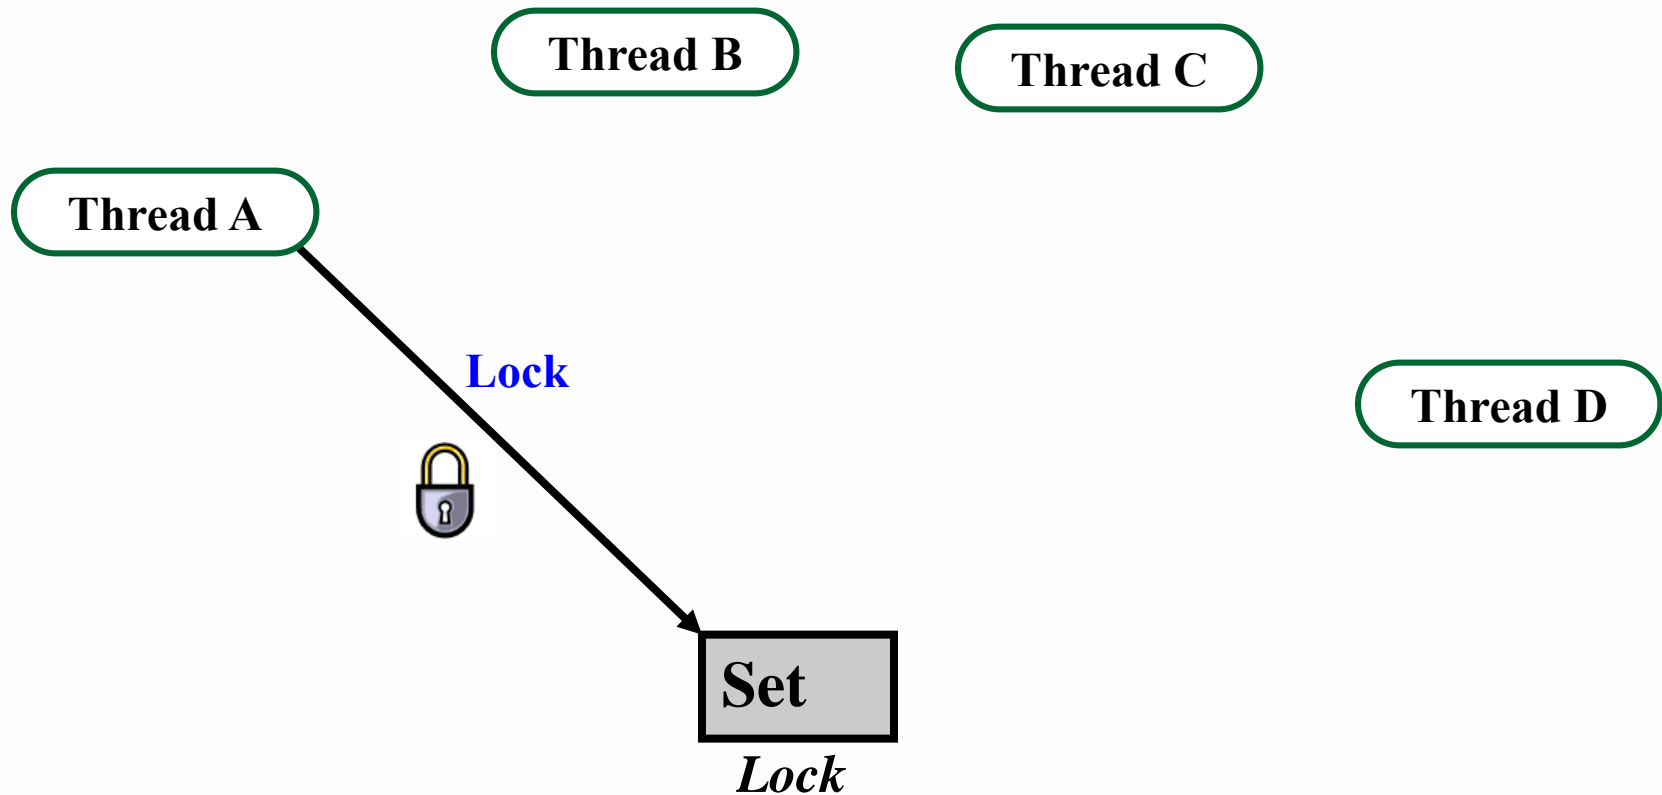
**Free**

*Lock*

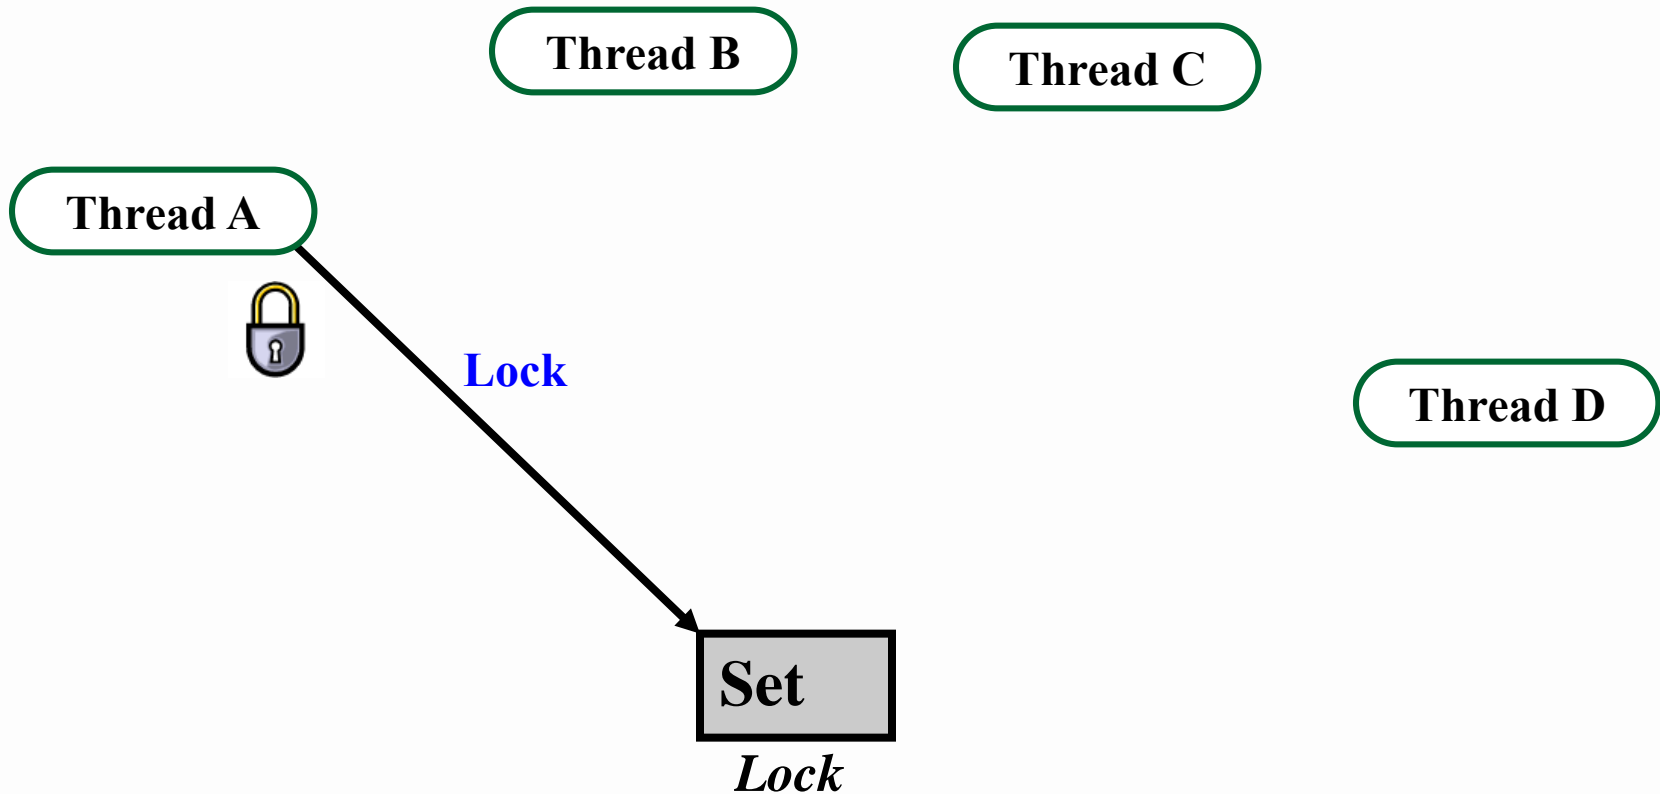
# Acquiring and Releasing Locks



# Acquiring and Releasing Locks



# Acquiring and Releasing Locks



# Acquiring and Releasing Locks

Thread B

Thread C

Thread A

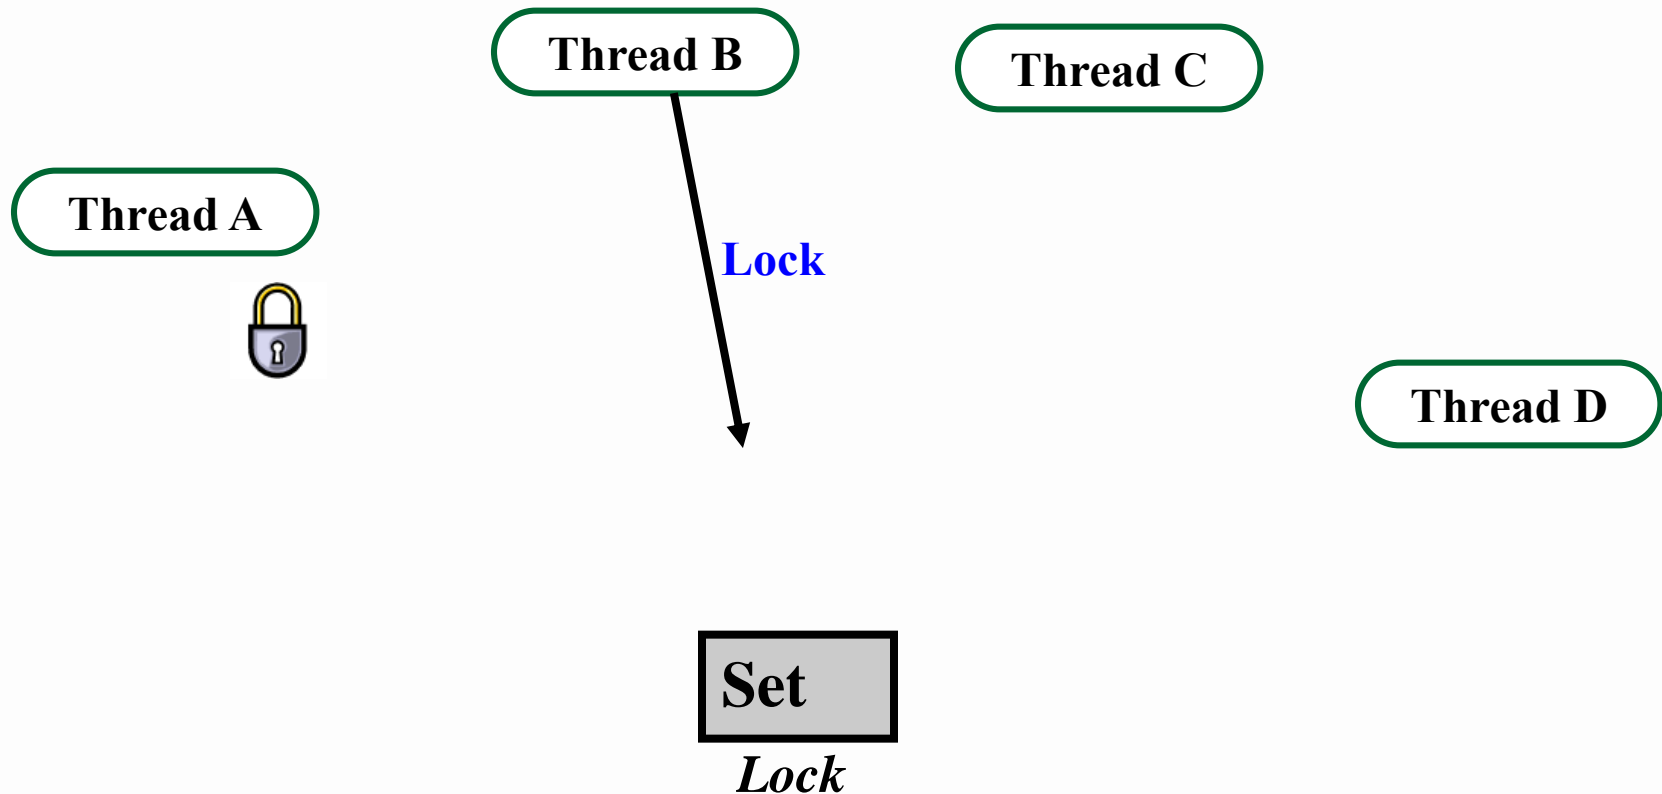


Thread D

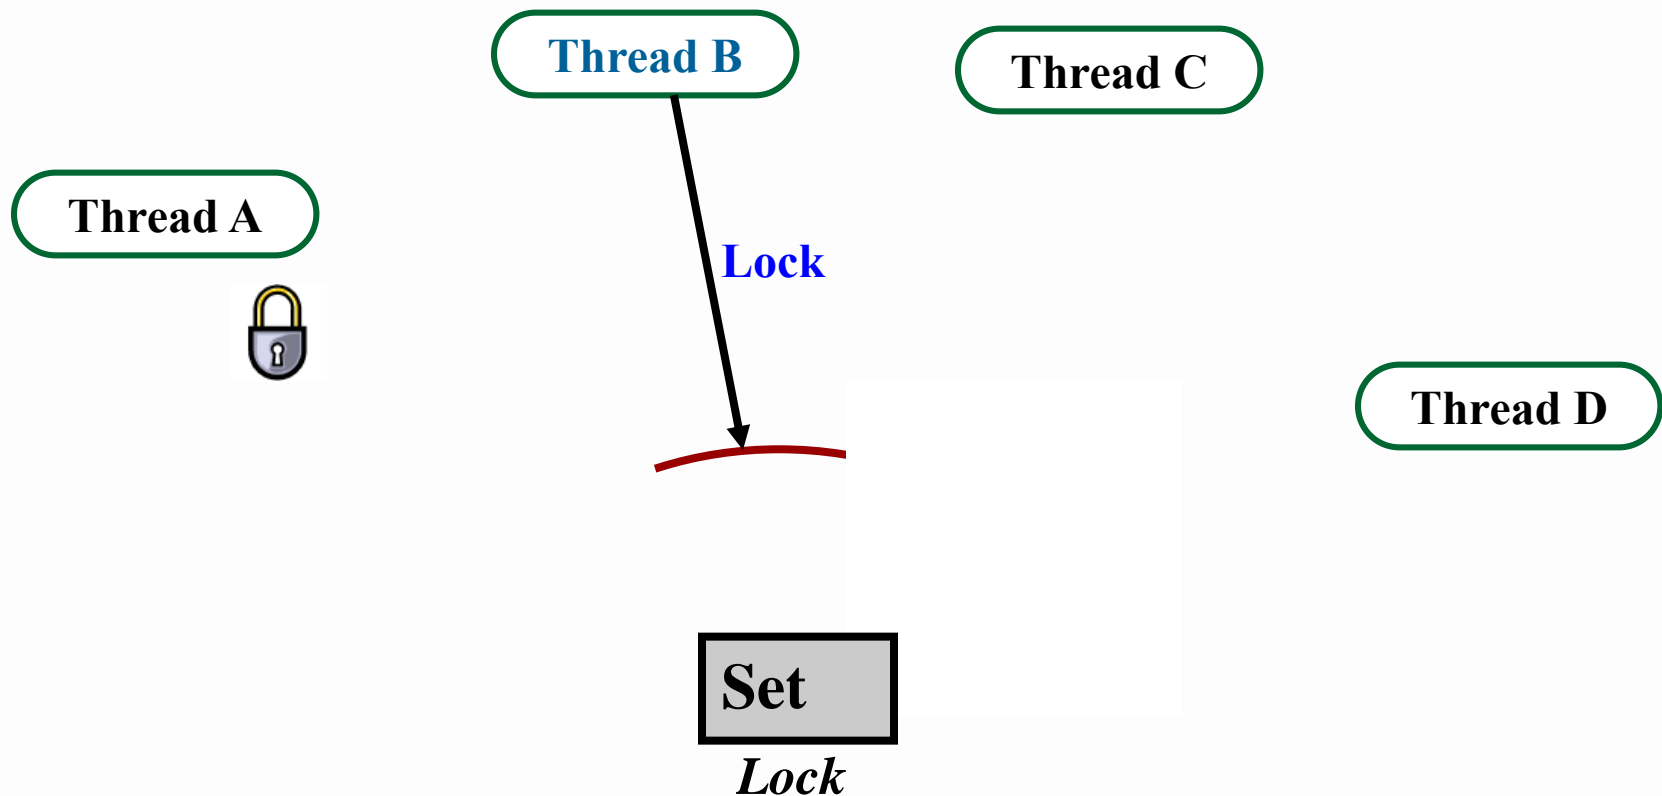
Set

*Lock*

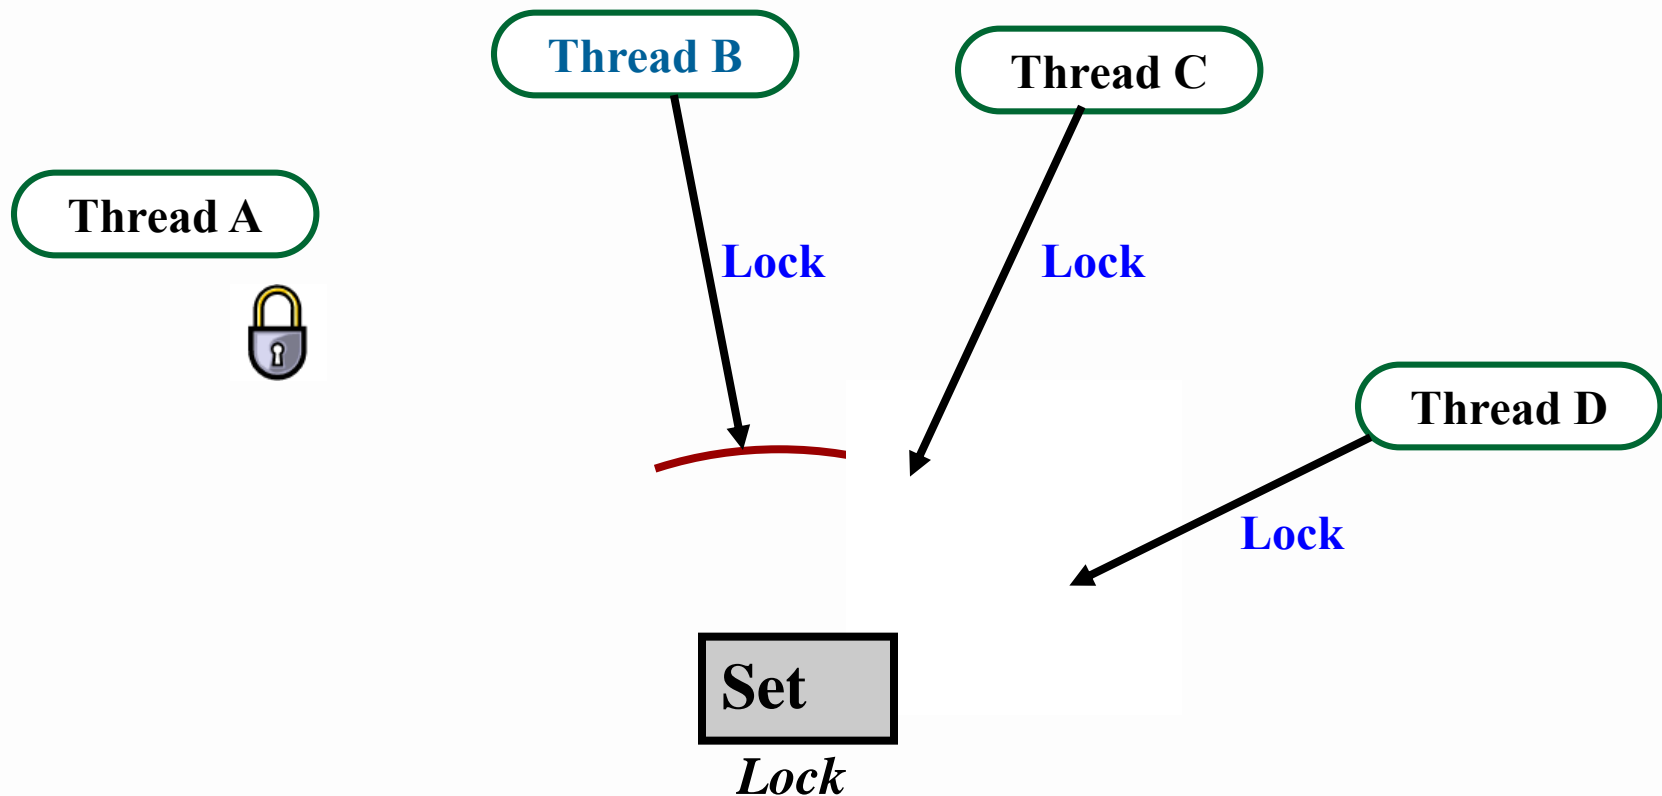
# Acquiring and Releasing Locks



# Acquiring and Releasing Locks

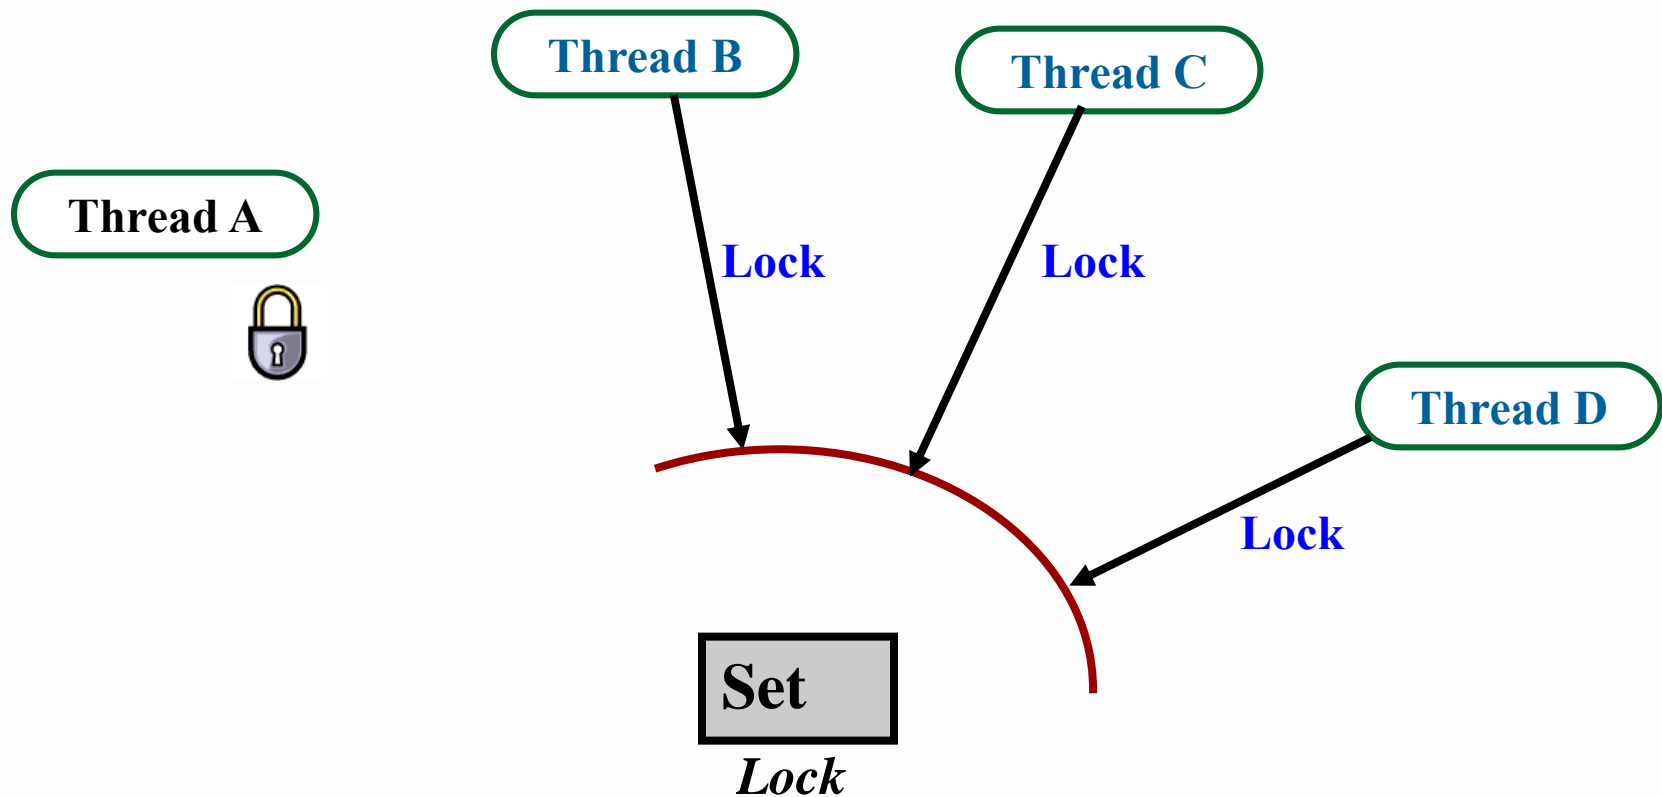


# Acquiring and Releasing Locks

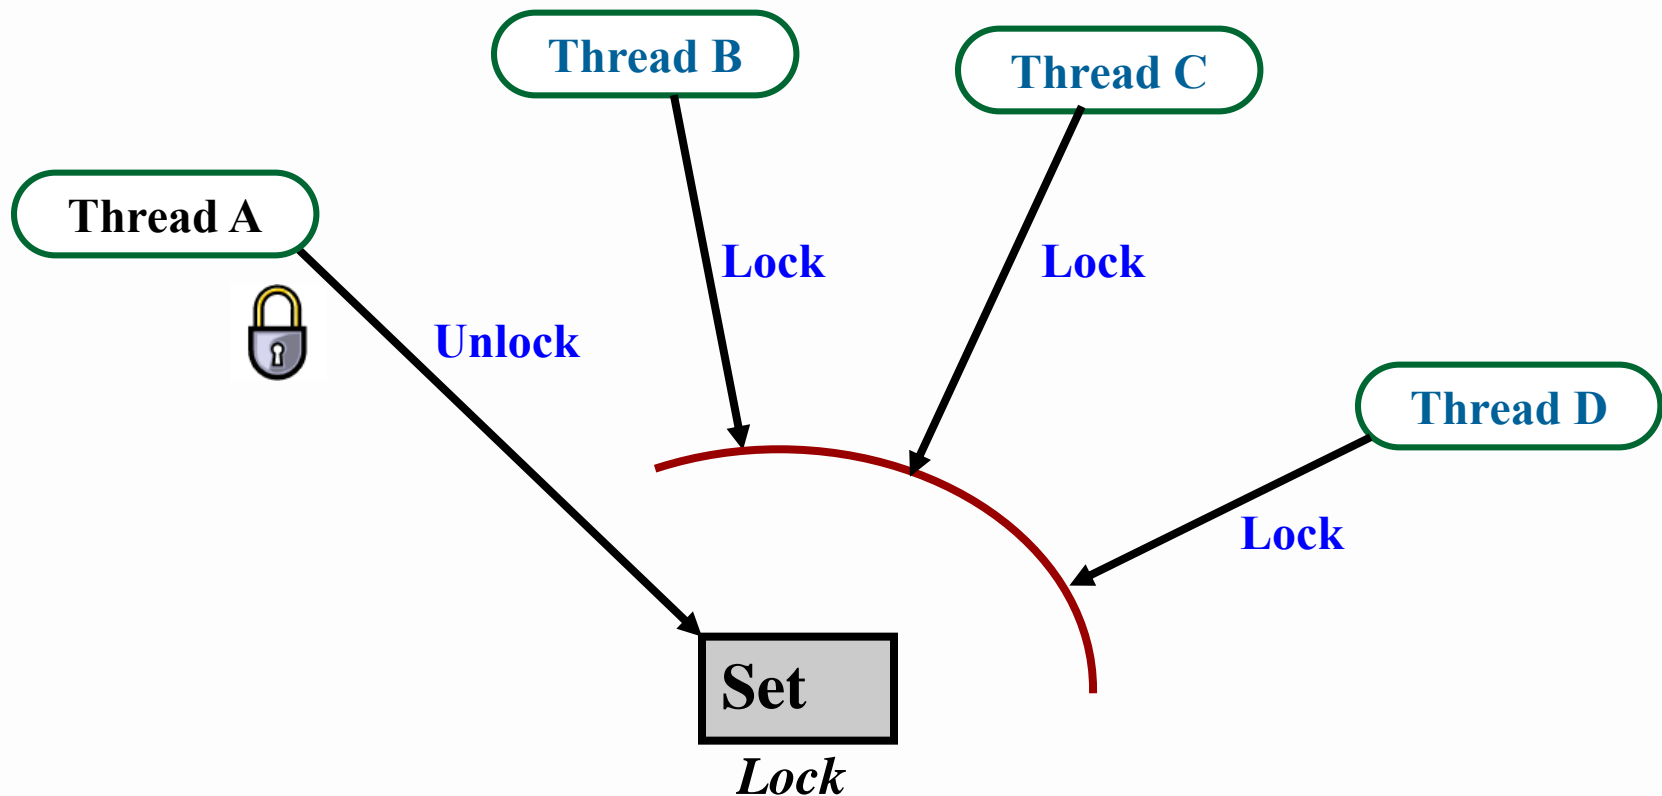




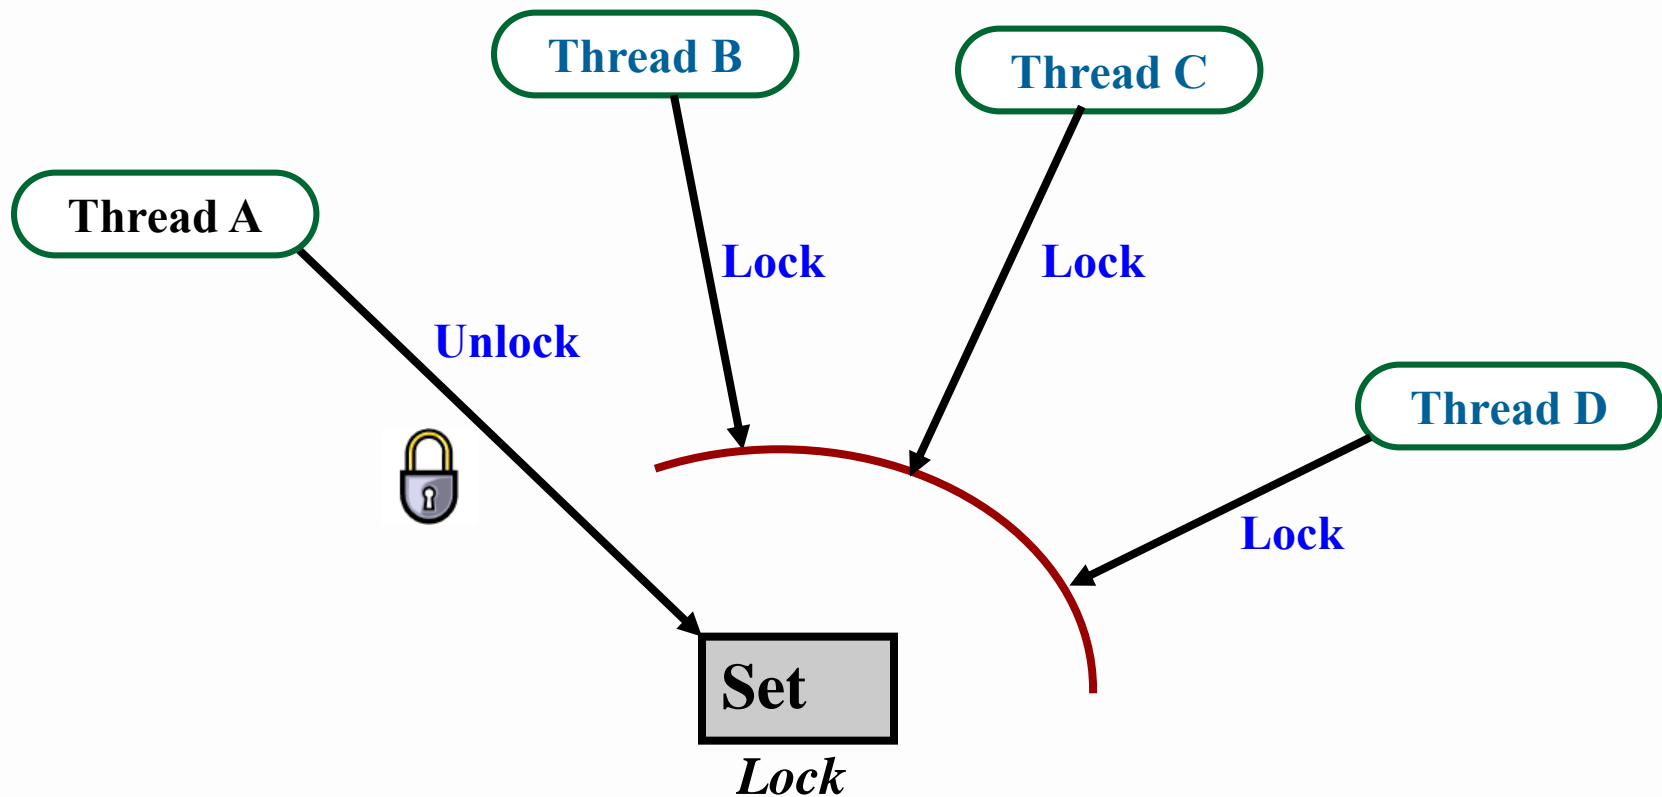
# Acquiring and Releasing Locks



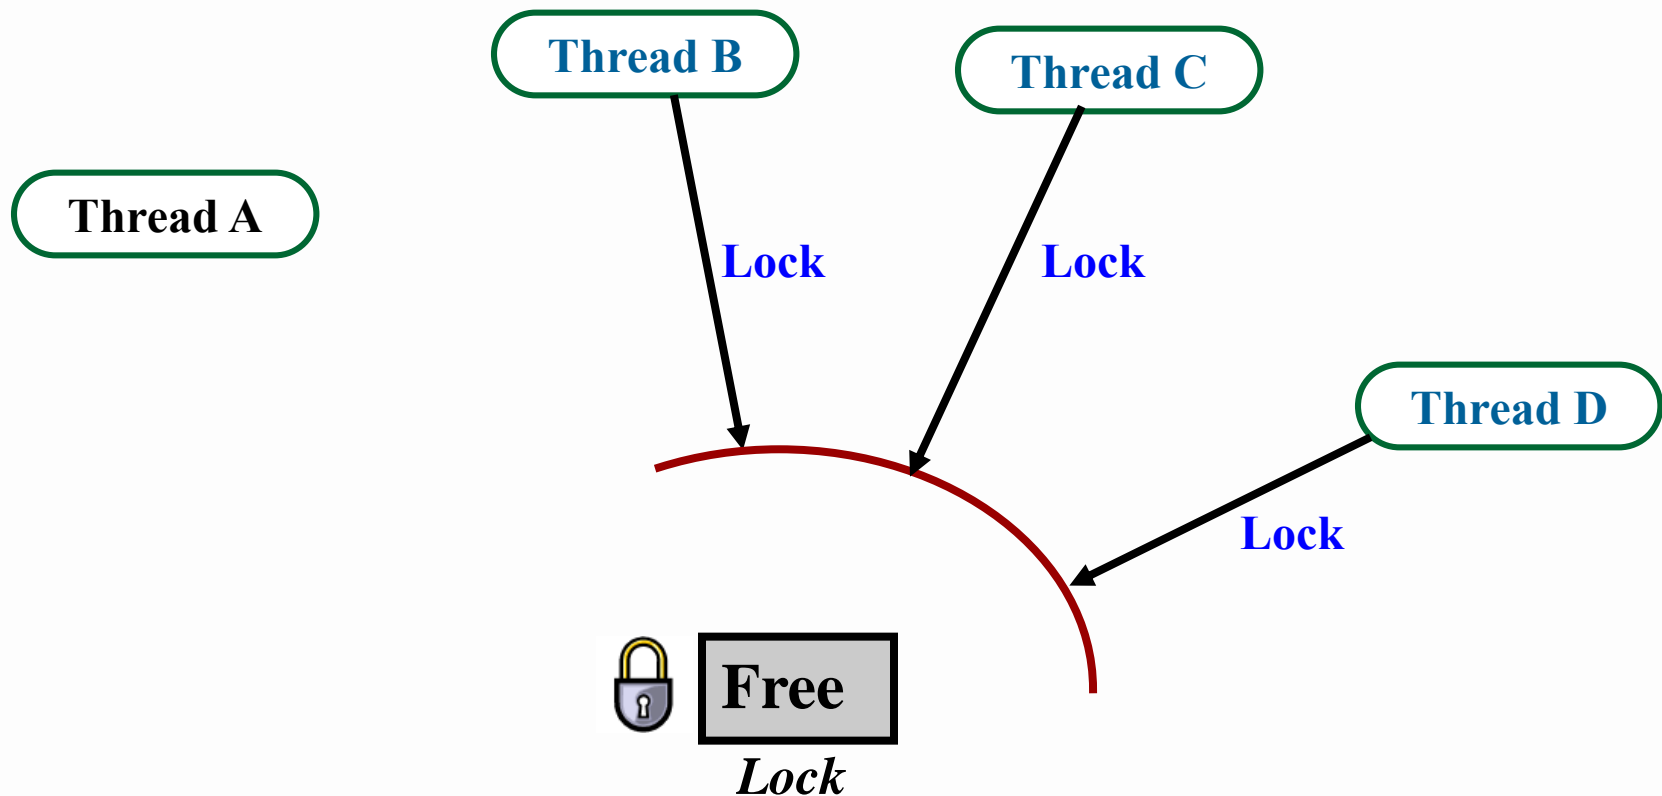
# Acquiring and Releasing Locks



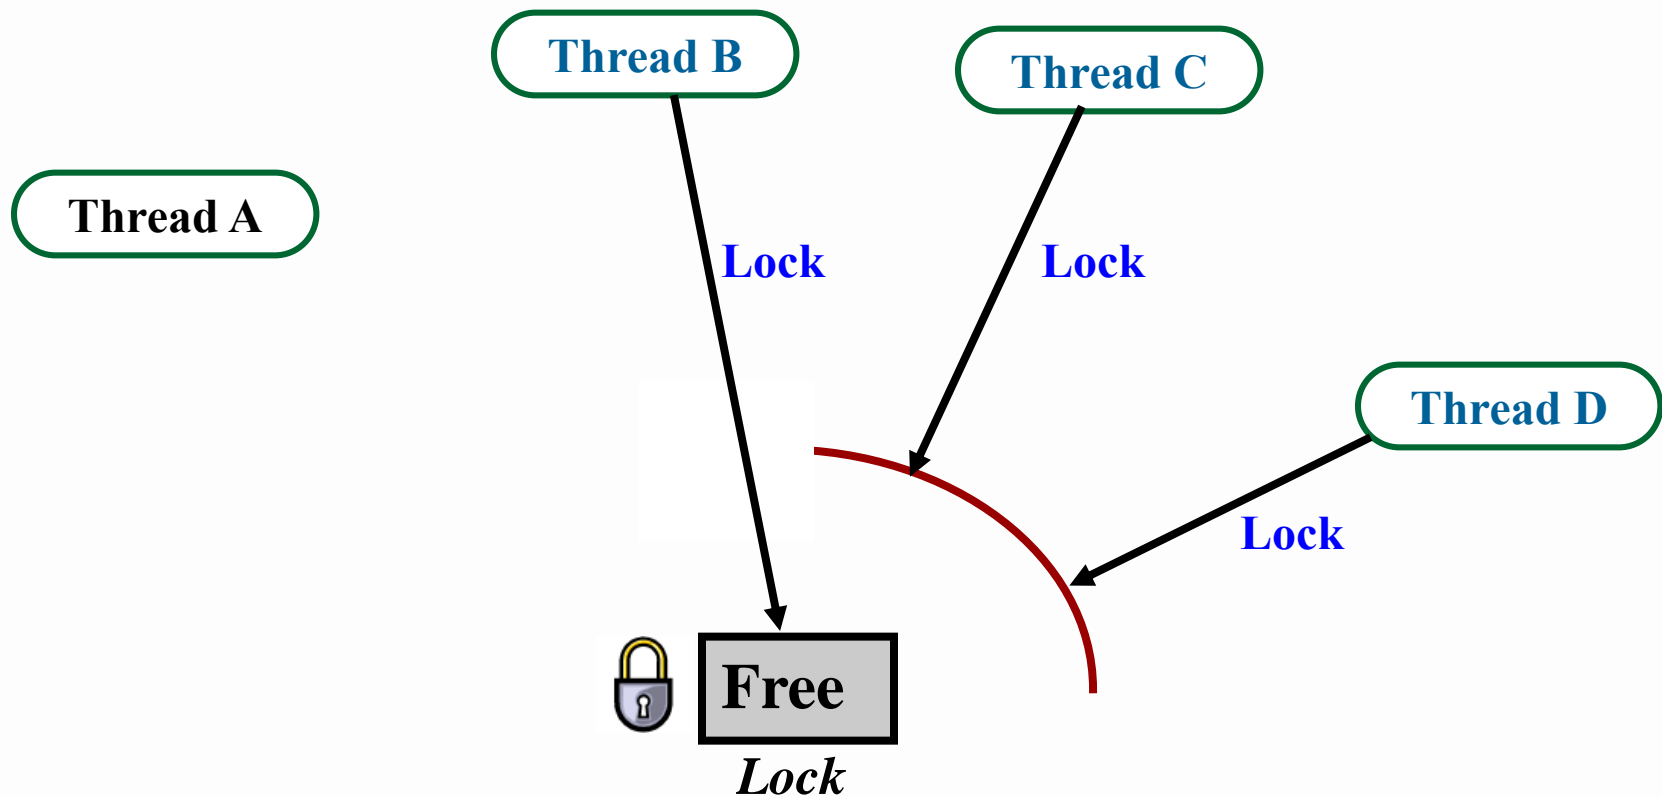
# Acquiring and Releasing Locks



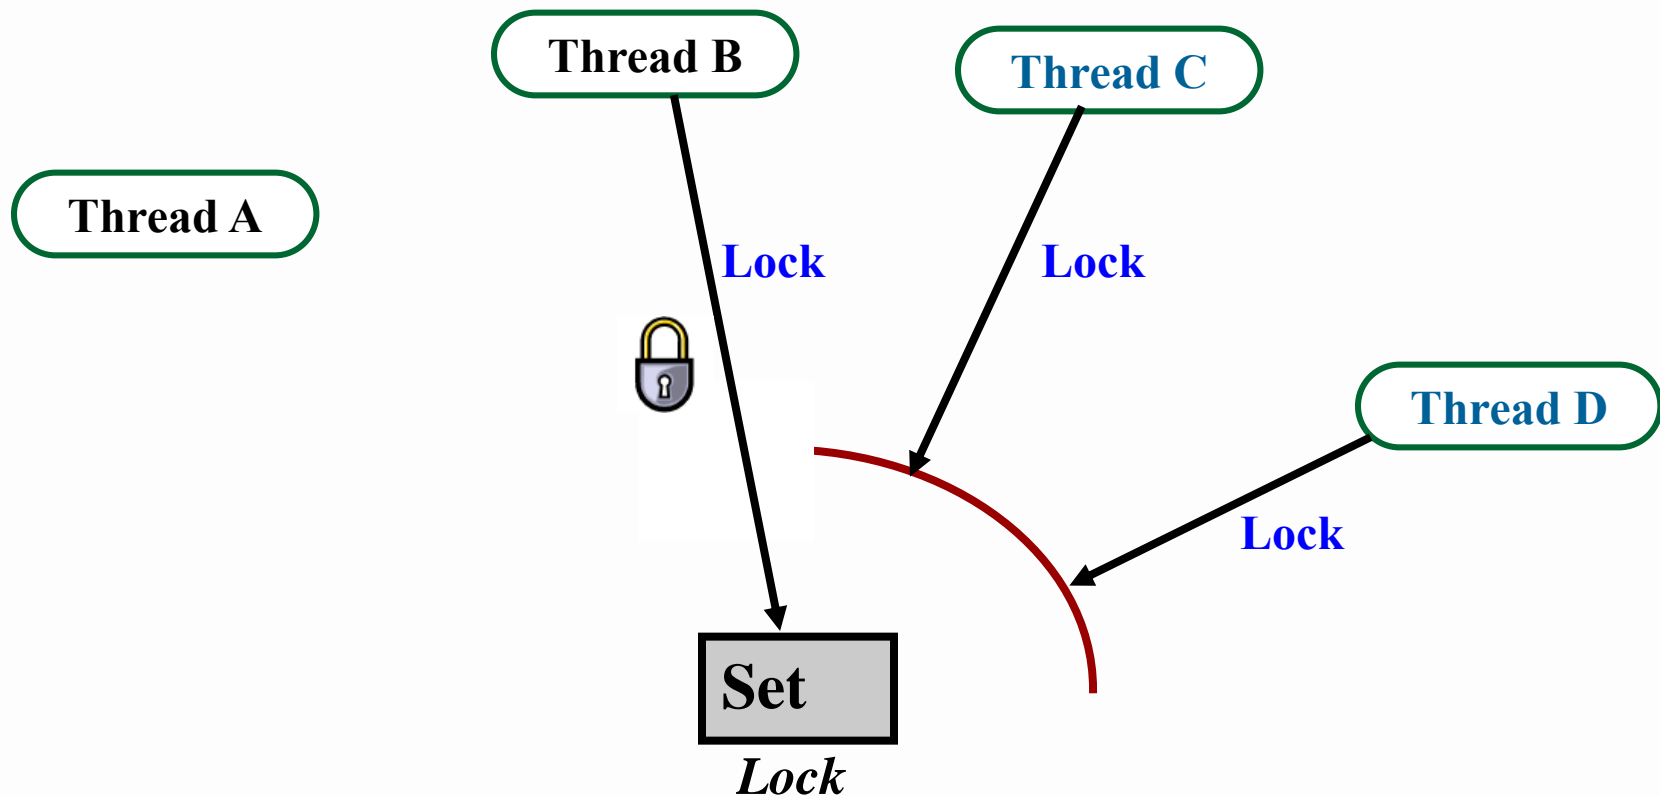
# Acquiring and Releasing Locks



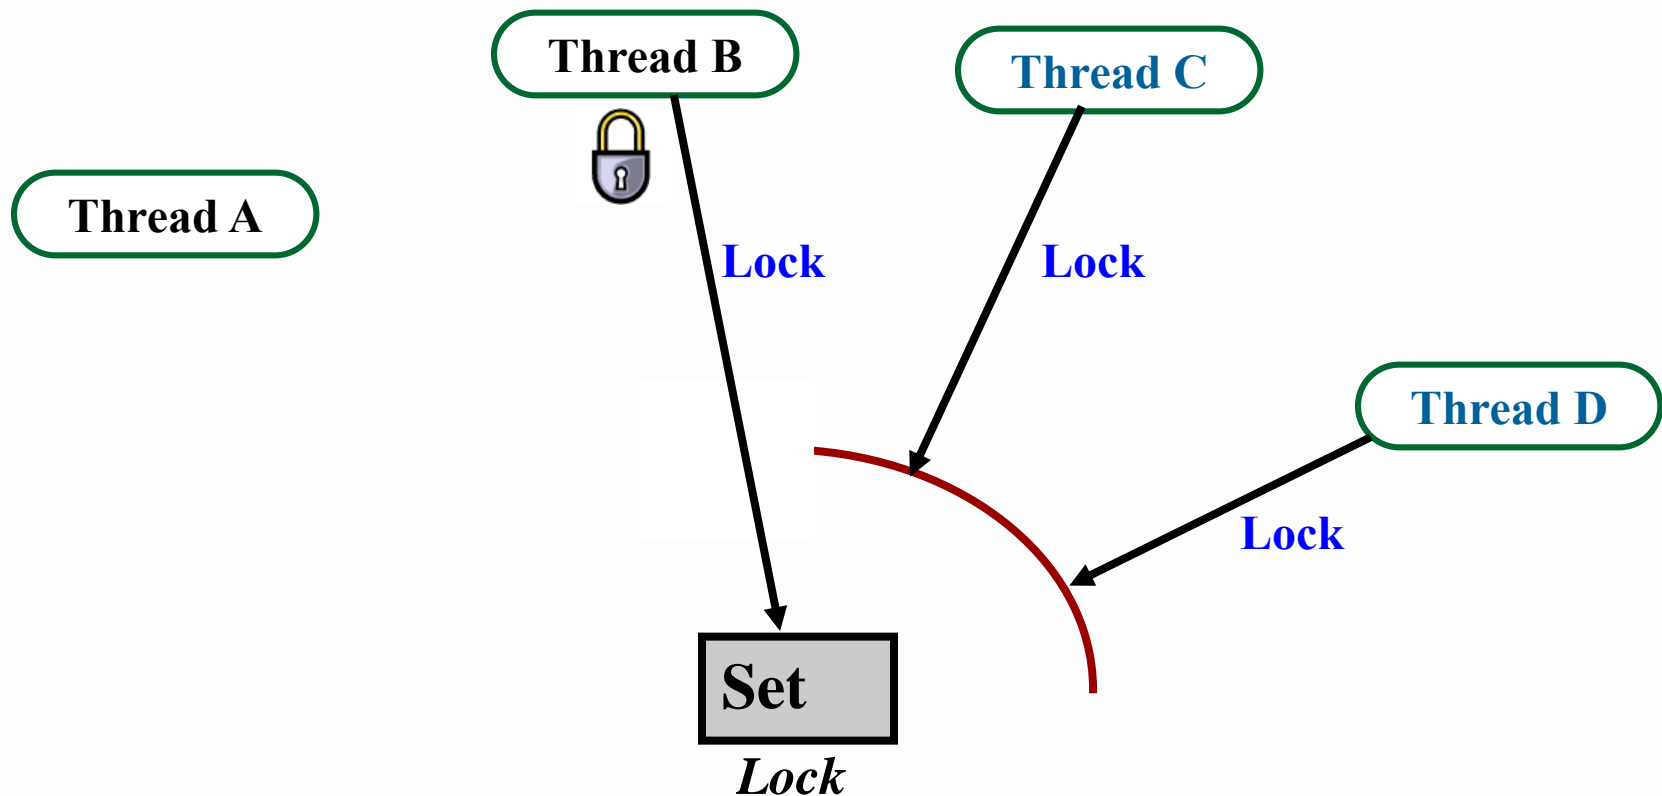
# Acquiring and Releasing Locks



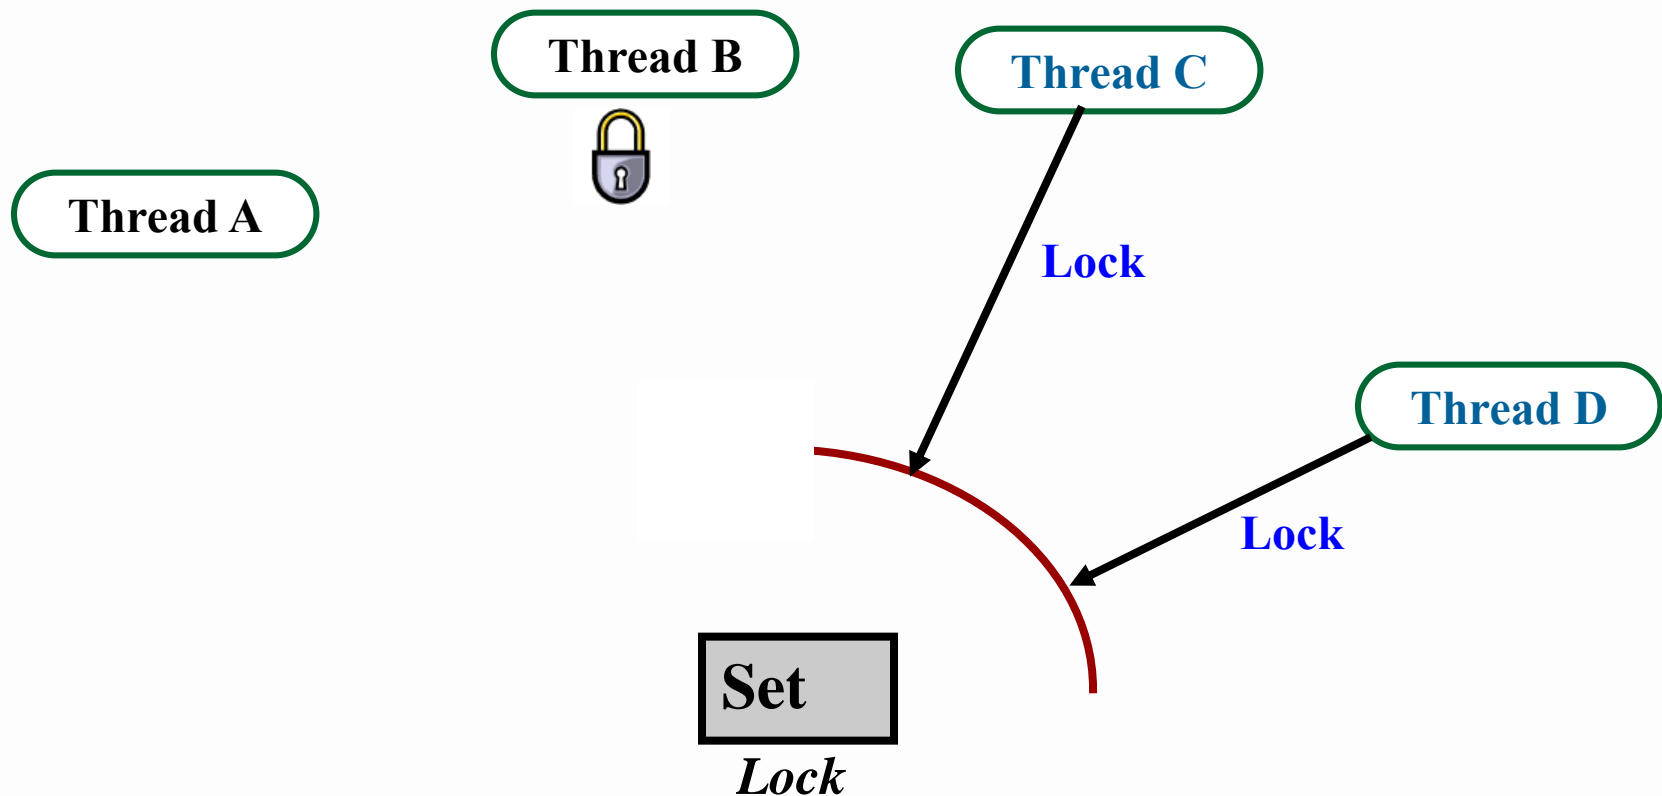
# Acquiring and Releasing Locks



# Acquiring and Releasing Locks



# Acquiring and Releasing Locks





# C implementation of Lock

```
#include <pthread.h>
```

```
1 pthread_mutex_t lock =  
  PTHREAD_MUTEX_INITIALIZER;
```

```
2 Pthread_mutex_lock(&lock);
```

```
3 balance = balance + 1;  
  // critical section
```

```
4 Pthread_mutex_unlock(&lock);
```

# Java implementation of Lock

- Using Lock object
- Using “Synchronized” keyword
  - synchronized methods
  - synchronized statements

# Locks (Interface)

```
public interface Lock {
```

```
    public void lock();
```

**acquire lock**

```
    public void unlock();
```

**release lock**

```
}
```

# Locks (Interface)

- Lock: interface
- Classes implement interface Lock
  - ReentrantLock
  - ReentrantReadWriteLock.ReadLock
  - ReentrantReadWriteLock.WriteLock

# Using ReentrantLock Class

```
public class MyClass {  
    private int shared_variable;  
    private ReentrantLock lock;  
    //ReentrantLock class implements Lock interface  
    public void myFunction() {  
        lock.lock();  
        try {  
  
            shared_variable++;  
  
        } finally {  
            lock.unlock();  
        }  
    }  
}
```

# Using Locks

```
public class MyClass {  
    private int shared_variable;  
    private ReentrantLock lock;  
    public void myFunction() {  
        lock.lock();  
        try {  
  
            shared_variable++;  
  
        } finally {  
            lock.unlock();  
        }  
    }  
}
```

**acquire Lock**

# Using ReentrantLock Class

```
public class MyClass {  
    private int shared_variable;  
    private ReentrantLock lock;  
    public void myFunction() {  
        lock.lock();  
        try {  
  
            shared_variable++;  
  
        } finally {  
            lock.unlock();  
        }  
    }  
}
```

Release lock  
(no matter what)

# Using ReentrantLock Class

```
public class MyClass {  
    private int shared_variable;  
    private ReentrantLock lock;  
    public void myFunction() {  
        lock.lock();  
        try {  
            shared_variable++;  
        } finally {  
            lock.unlock();  
        }  
    }  
}
```

critical section



# Java implementation of Lock

- Using Lock object
- Using “Synchronized” keyword
  - synchronized methods
  - synchronized statements

# Using “Synchronized” keyword

- Each Java object has an associated intrinsic lock
  - The lock is initially un-owned
  - As long as a thread owns an intrinsic lock, no other thread can acquire the same lock. The other thread will block when it attempts to acquire the lock.
- **synchronized** keyword forces the current thread to obtain an object's intrinsic lock

# ME in Java using **synchronized**

```
public synchronized void MyFunction() {  
  
    // critical section  
  
}
```



```
synchronized(shared_object) {  
  
    // critical section  
  
}
```

# Example: Synchronized methods

```
public class SynchronizedCounter {  
  
    private int c = 0;  
  
    public synchronized void increment() {  
        c++;  
    }  
  
    public synchronized void decrement() {  
        c--;  
    }  
  
    public synchronized int value() {  
        return c;  
    }  
}
```

# Two effects of Synchronized methods

- First, it is not possible for two invocations of synchronized methods on the same object to **interleave**. When one thread is executing a synchronized method for an object, all other threads that invoke synchronized methods for the same object block (suspend execution) until the first thread is done with the object.
- Second, when a synchronized method exits, it automatically establishes a happens-before relationship with any subsequent invocation of a synchronized method for the same object. **This guarantees that changes to the state of the object are visible to all threads.**

# Synchronized methods

- Question:
  - Can constructors be synchronized?
- Answer:
  - No
  - using the synchronized keyword with a constructor is a syntax error. Synchronizing constructors doesn't make sense, because only the thread that creates the object should have access to it while the object is being constructed

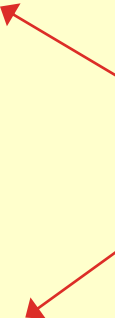
# Example: Synchronized statements

```
public void addName(String name) {  
  
    // other work goes here  
  
    synchronized(this) {  
  
        lastName = name;  
        nameCount++;  
  
    }  
  
    // other work goes here  
  
}
```

# Example: Synchronized statements

```
public class MsLunch {  
    private long c1 = 0;  
    private long c2 = 0;  
    private Object lock1 = new Object();  
    private Object lock2 = new Object();  
  
    public void inc1() {  
        synchronized(lock1) {  
            c1++;  
        }  
    }  
  
    public void inc2() {  
        synchronized(lock2) {  
            c2++;  
        }  
    }  
}
```

Using an object that  
provides intrinsic lock





# These two are equivalent

```
public synchronized void MyFunction() {  
  
    // critical section  
  
}
```

```
public void MyFunction() {  
    synchronized(this) {  
  
        // critical section  
  
    }  
}
```

# Using synchronized

- The synchronized keyword automatically implements:
  - lock
  - try
  - finally
  - unlock

# Which one is better?

## Lock interface or synchronization blocks?

- **Answer:** Lock interface (Reentrant locks)!
- Lock interface has more benefits...

# Which one is better?

## Lock interface or synchronization blocks?

- With the Lock interface, you have the option of calling ***tryLock()***

```
if(lock.tryLock() ) {  
    //Lock is acquired  
} else{  
    //Lock isn't acquired; act accordingly  
}
```

- You also have the option to specify a timeout (from nanoseconds to days!)

```
if(lock.tryLock(50, TimeUnit.SECONDS) ) {  
    //Lock is acquired  
} else{  
    //Lock isn't acquired; act accordingly  
}
```

# Which one is better?

## Lock interface or synchronization blocks?

- **Answer: Lock interface (Reentrant locks)!**
- A thread blocked on an intrinsic lock cannot be interrupted. All or nothing. (inflexible)
- If a thread becomes blocked and the semantics of your program will never allow it to become unblocked (it's stuck until you kill the JVM)
- With the Lock interface, you use it with the ***lockInterruptibly()*** method

# ReentrantLock interface

```
public interface Lock {  
  
    public void lock();  
  
    public void lockInterruptibly();  
  
    public void unlock();  
  
    public Boolean tryLock();  
  
    public Boolean tryLock(long time, TimeUnit  
unit);  
  
}
```

# Acknowledgement

- Chapter 26
  - Operating Systems: Three Easy Pieces
- Java documentation
  - <https://docs.oracle.com/javase/tutorial/essential/concurrency/index.html>
- 4.ppt
  - Intro to Operating System at Portland State University
  - by Jonathan Walpole

# Questions?