



COS40003 Concurrent Programming

Lecture 8: Semaphore

Outline

- Why semaphores?
- What is a semaphore?
- How to use semaphores?
 - Binary semaphores (similar to locks)
 - Semaphore for ordering (similar to condition variable)
 - Producer and Consumer using semaphore
 - Implementing semaphore with lock and condition variable
- Java methods


Why semaphore?

- We introduced **Locks** and **Condition Variables**, and we can write good concurrent programs using them.
- Why do we need semaphore?
- Reason
 - Two concepts share similar idea, but not the same thing. They were proposed by different people.

Why semaphore?


- Semaphore

- 1968 by Dijkstra
- Proposed “private semaphore” in “Cooperating sequential processes”



- Condition variable

- 1974 by Hoare
- In his work on “monitors” in “Monitors: An Operating System Structuring Concept”



Outline

- Why semaphores?
- What is a semaphore?
- How to use semaphores?
 - Binary semaphores (similar to locks)
 - Semaphore for ordering (similar to condition variable)
 - Producer and Consumer using semaphore
 - Implementing semaphore with lock and condition variable
- Java methods

What is a semaphore?

- A semaphore is an object with an integer value that we can manipulate with two routines;
 - P/down/wait/ acquire (decrements the count)
 - V/ up /post/ release (increments the count)
 - Eg., P(semaphore), V(semaphore)
- P,V – (from Dutch words prolaag, verlaag)
- wait, post – POSIX standard
- acquire, release – Java API

Initialize a semaphore

```
#include <semaphore.h>
sem_t s;
sem_init(&s, 0, 1);
```

- 1. declare a semaphore s
- 2. second argument of sem_init() set to 0, this indicates the semaphore is shared between threads in the same process. (A different value for sharing across different processes.)
- 3. third argument, initialize the semaphore's value

P (wait) and V (post)

- sem_wait() - atomic
- sem_post() - atomic

```
int sem_wait(sem_t *s) {  
    decrement the value of semaphore s by one  
    wait if value of semaphore s is negative  
}  
  
int sem_post(sem_t *s) {  
    increment the value of semaphore s by one  
    if there are one or more threads waiting, wake one  
}
```


Discussion of wait and post

- 1. `sem_wait()`
 - either return right way,
 - Or cause the caller to suspend, waiting for a subsequent post.
- 2. `sem_post()`
 - simply increments the value of the semaphore and then, if there is a thread waiting to be woken, wakes one of them up.
- 3. The value of the semaphore, when negative, its additive inverse is equal to the number of waiting threads

Break

- Ignore how `sem_wait()` and `sem_post()` are implemented atomically, which will be discussed later
- First see how semaphore can be used?

Example: hot desks

Semaphore
value

2

available

2

entrance



desk please!

desk 1



desk 2



desk 3



1

available

1

entrance

desk 1



desk 2



desk 3



1

available

1

entrance



desk please!

desk 1



desk 2



desk 3



0

available

0

entrance



desk please!

desk 1



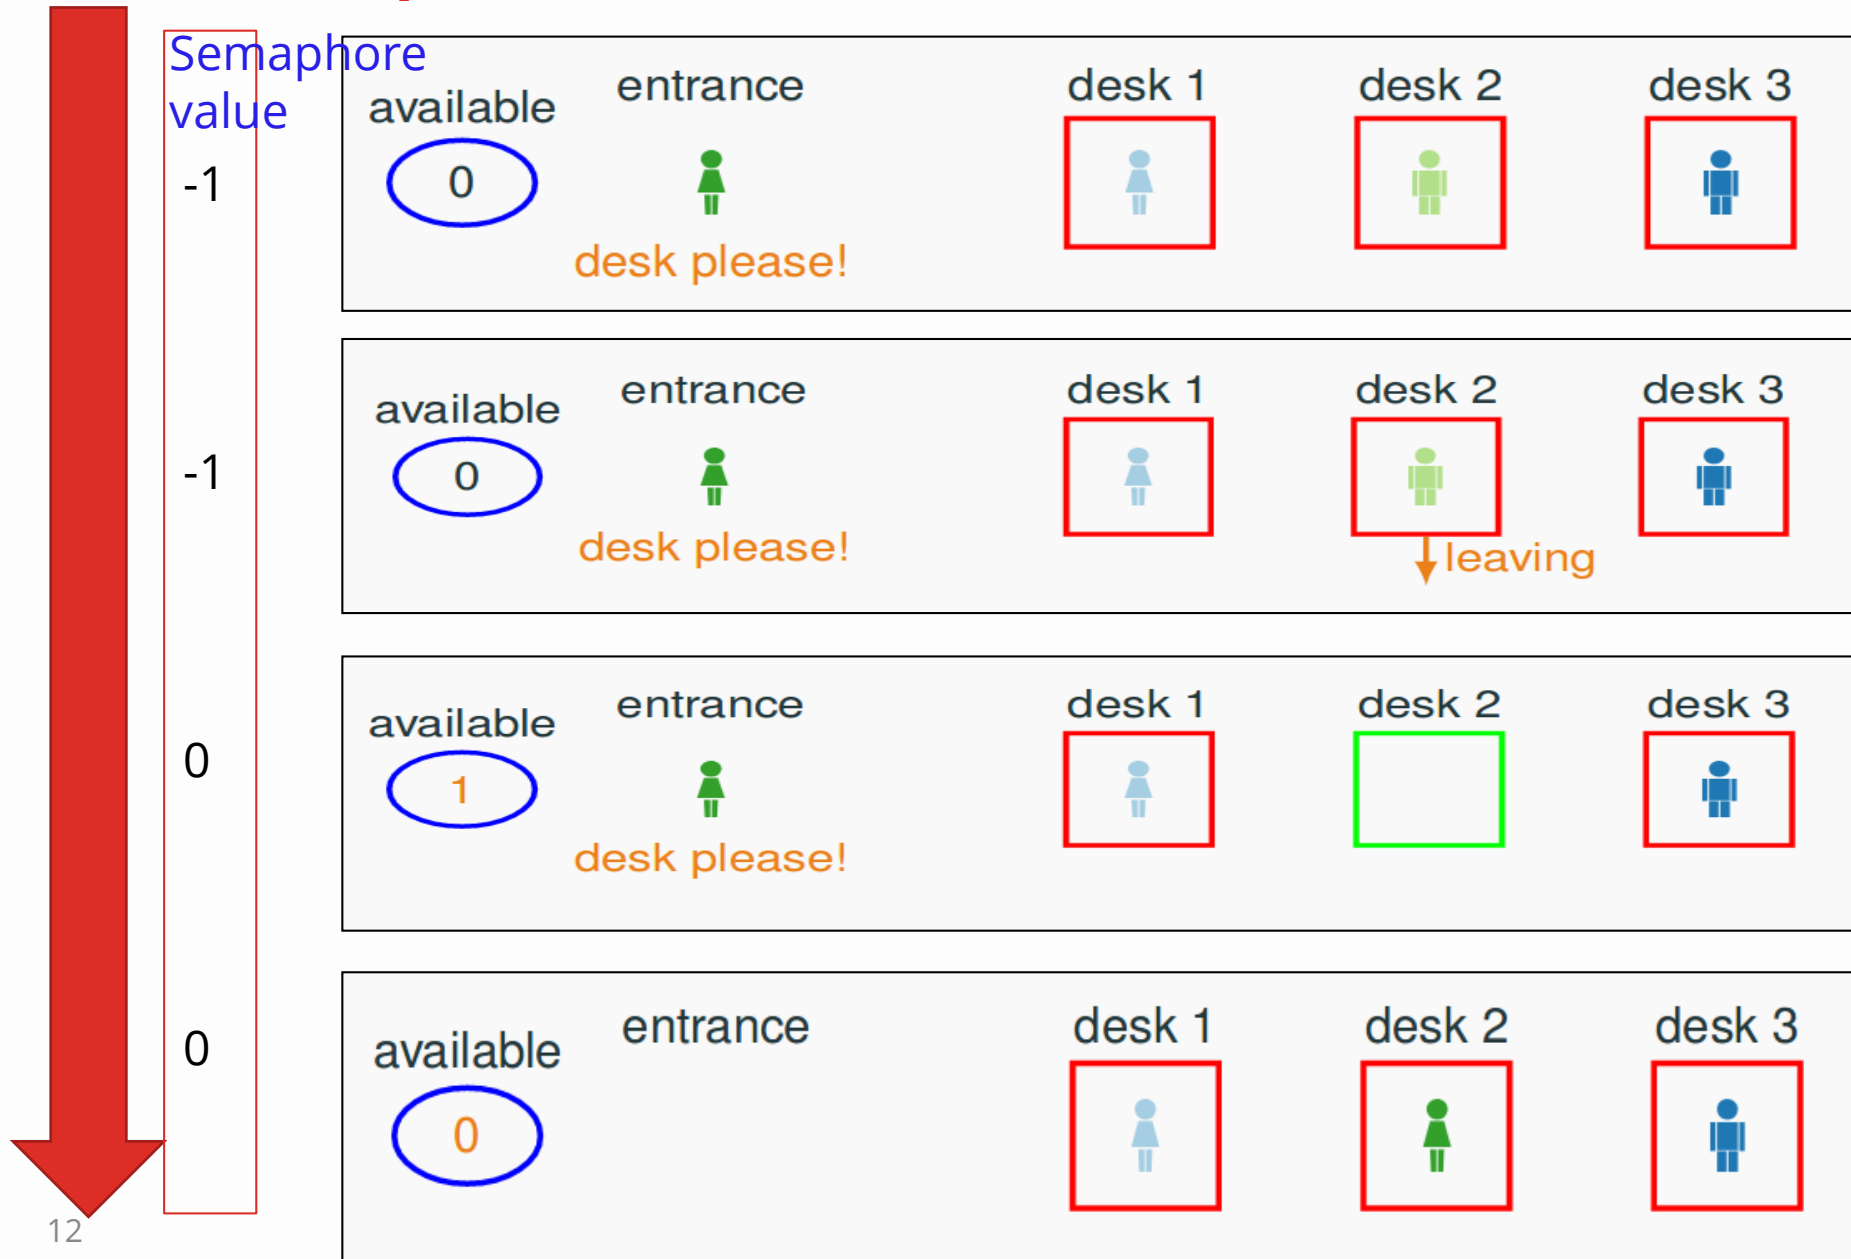
desk 2



desk 3



Example: hot desks



Outline

- Why semaphores?
- What is a semaphore?
- How to use semaphores?
 - Binary semaphores (similar to locks)
 - Semaphore for ordering (similar to condition variable)
 - Producer and Consumer using semaphore
 - Implementing semaphore with lock and condition variable
- Java methods

Binary Semaphores (Locks)

- Using a semaphore as a lock

```
sem_t m;  
sem_init(&m, 0, X);
```

```
// initialize semaphore to X; what should X be?
```

```
sem_wait(&m);  
    // critical section here  
sem_post(&m);
```

- $X=1$

Binary Semaphores (Locks): running example 1

- Case 1:
 - Thread 0 calls `sem_wait()`;
 - Semaphore value $X=1$ decreases to $X=0$;
 - $X \geq 0$, so return from `sem_wait()`;
 - Thread 0 enters critical section and finishes;
 - Thread 0 calls `sem_post()`;
 - Semaphore value $X=0$ increases to $X=1$;
 - Wake any waiting thread, none in this case, we are done.

Binary Semaphores (Locks): running example 2

- Case 2:

- Thread 0 calls `sem_wait()`;
- Semaphore value $X=1$ decreases to $X=0$;
- $X \geq 0$, so return from `sem_wait()`;
- Thread 0 enters critical section;

--context switch, recall only `wait()`, `post()` are atomic--

- Thread 1 calls `sem_wait()`;
- Semaphore value $X=0$ decreases to $X= -1$;
- $X < 0$, so Thread 1 puts itself to sleep, waiting;
- (To be continued)

Binary Semaphores (Locks): running example 2

- Case 2:
 - (continue with the previous page)
--context switch to Thread 0, because Thread 1 is sleeping--
 - Thread 0 finishes and calls `sem_post()`;
 - Semaphore value $X = -1$ increases to $X = 0$;
 - Wake any waiting thread, in this case, Thread 1 waken up.
 - Thread1 enters critical section and finishes
 - Thread 1 calls `sem_post()`;
 - Semaphore value $X = 0$ increases to $X = 1$;
 - Wake any waiting thread, in this case, none and we are done

Binary Semaphores (Locks): conclusion

- Binary semaphores can simulate locks
- Question: any difference you find between binary semaphores and locks?
- Answer: semaphore support transferring of permissions
 - using a lock, only the thread decrements the counter can increment it back to 1;
 - using a semaphore, one thread decrements the counter to 0 and may let another thread increment it to 1.

Outline

- Why semaphores?
- What is a semaphore?
- How to use semaphores?
 - Binary semaphores (similar to locks)
 - Semaphore for ordering (similar to condition variable)
 - Producer and Consumer using semaphore
 - Implementing semaphore with lock and condition variable
- Java methods

Recall the example: parent waiting for child

- 1 void *child(void *arg) {
2 printf("child\n");
3 // XXX how to indicate we are done?
4 return NULL;
5 }
6
7 int main(int argc, char *argv[]) {
8 printf("parent: begin\n");
9 pthread_t c;
10 Pthread_create(&c, NULL, child, NULL);
11 // create child
11 // XXX how to wait for child?
12 printf("parent: end\n");
13 return 0;
14 }

Semaphores for Ordering (similar to condition variable)

```
sem_t s;  
void * child(void *arg) {  
    printf("child\n");  
    sem_post(&s); // signal here: child is done  
    return NULL;  
}  
  
int main(int argc, char *argv[]) {  
    sem_init(&s, 0, X); // what should X be?  
    printf("parent: begin\n");  
    pthread_t c;  
    Pthread_create(&c, NULL, child, NULL);  
    sem_wait(&s); // wait here for child  
    printf("parent: end\n");  
    return 0;  
}
```

Answer:

X=0

Semaphores For Ordering (similar to condition variable)

```
sem_t s;
void * child(void *arg) {
    printf("child\n");
    sem_post(&s); // signal here: child is done
    return NULL;
}

int main(int argc, char *argv[]) {
    sem_init(&s, 0, X); // what should X be?
    printf("parent: begin\n");
    pthread_t c;
    Pthread_create(&c, NULL, child, NULL);
    sem_wait(&s); // wait here for child
    printf("parent: end\n");
    return 0;
}
```

Case 1: parent runs before child

1. Parent calls wait() before child calls post();
2. Semaphore $X=0$ decreased to $X=-1$
3. $X<0$, so parent puts itself to sleep
4. Child calls post;
5. Semaphore $X=-1$ increased to $X=0$;
6. Parent waken up and finishes

Semaphores For Ordering (similar to condition variable)

```
sem_t s;
void * child(void *arg) {
    printf("child\n");
    sem_post(&s); // signal here: child is
    return NULL;
}

int main(int argc, char *argv[]) {
    sem_init(&s, 0, X); // what should X be
    printf("parent: begin\n");
    pthread_t c;
    Pthread_create(&c, NULL, child, NULL);
    sem_wait(&s); // wait here for child
    printf("parent: end\n");
    return 0;
}
```

Case 2: parent runs after child

1. child calls post() before parent calls wait();
2. Semaphore X=0 increased to X=1
3. Wake up any one waiting, in this case, none
4. Later, parent calls wait;
5. Semaphore X=1 decreased to X=0;
6. $X \geq 0$, parent will not be put to sleep. Parent returns from wait() and go on

Outline

- Why semaphores?
- What is a semaphore?
- How to use semaphores?
 - Binary semaphores (similar to locks)
 - Semaphore for ordering (similar to condition variable)
 - Producer and Consumer using semaphore
 - Implementing semaphore with lock and condition variable
- Java methods

Semaphore in the Producer/Consumer Problem

- Recall: In Lecture 7
- A consumer should only wake producers, and a producer should only wake consumers.
 - Solution: two condition variable
 - Otherwise: wrong wakeups

Semaphore in the PC Problem - First attempt

- Use two semaphores, empty and full
- The producer first waits for a buffer to become empty in order to put data into it, and the consumer similarly waits for a buffer to become filled before using it.

Semaphore in the PC Problem - First attempt

```
int buffer[MAX];  
int fill = 0;  // put pointer  
int use = 0;  // get pointer  
  
void put(int value) {  
    buffer[fill] = value;  
    fill = (fill + 1) % MAX;  
}  
  
int get() {  
    int tmp = buffer[use];  
    use = (use + 1) % MAX;  
    return tmp;  
}
```

```
sem_t empty; sem_t full;
void *producer(void *arg) {
    for (int i = 0; i < loops; i++) {
        sem_wait(&empty);
        put(i);
        sem_post(&full);
    }
}
```

```
void *consumer(void *arg) {
    while (tmp != -1) {
        sem_wait(&full);
        tmp = get();
        sem_post(&empty);
    }
}
```

```
int main(int argc, char *argv[]) {
    sem_init(&empty, 0, MAX); // MAX buffers are empty to begin
    sem_init(&full, 0, 0); // ... and 0 are full
    ...
}
```

Is it correct?

Suppose
MAX=1, let us
examine

```

sem_t empty; sem_t full;
void *producer(void *arg) {
    for (int i = 0; i < loops; i++) {
        sem_wait(&empty);
        put(i);
        sem_post(&full);
    }
}

void *consumer(void *arg) {
    while (tmp != -1) {
        sem_wait(&full);
        tmp = get();
        sem_post(&empty);
    }
}

int main(int argc, char *argv[]) {
    sem_init(&empty, 0, MAX); // MAX buffers are empty
    sem_init(&full, 0, 0); // ... and 0 are full
    ...
}

```

Example: MAX = 1

1. A consumer runs first, calls wait(&full), full is set to -1, consumer is blocked waiting for a post(&full);
2. A producer comes, calls wait(&empty), empty is set to 0 from MAX=1.
3. The producer puts data into the buffer and calls post(&full), full will be set to 0 from -1, wakes up the consumer.
4. (1) The producer keeps running, calls wait() again, empty is set to -1, the producer is blocked;
5. (2) The consumer waken up, calls get(), then post(&empty)

Semaphore in the PC Problem - First attempt

- Discussion:
- Question, suppose $MAX = 1$, will it work with “multiple producers, and multiple consumers”?
- Answer: Yes

Semaphore in the PC Problem - First attempt

- Discussion:
- Question 2, suppose $MAX=10$, will it work with “multiple producers, and multiple consumers”?
- Answer: No

```

sem_t empty; sem_t full;
void *producer(void *arg) {
    for (int i = 0; i < loops; i++) {
        sem_wait(&empty);
        put(i);
        sem_post(&full);
    }
}

void *consumer(void *arg) {
    while (tmp != -1) {
        sem_wait(&full);
        tmp = get();
        sem_post(&empty);
    }
}

int main(int argc, char *argv[]) {
    sem_init(&empty, 0, MAX); // MAX buffers are empty to begin
    sem_init(&full, 0, 0); // ... and 0 are full
    ...
}

```

Recall:

only wait() and post(),
the functions
themselves are
guaranteed to be
atomic!

Semaphore in the PC Problem - First attempt

```
int buffer[MAX];
int fill = 0;  // put pointer
int use = 0;  // get pointer

void put(int value) {
    buffer[fill] = value;
    fill = (fill + 1) % MAX;
}

int get() {
    int tmp = buffer[use];
    use = (use + 1) % MAX;
    return tmp;
}
```

Thread 1 calls put();
Thread 2 call put();

Sequence is:

Thread1:
buffer[fill] = value;

Thread 2:
buffer[fill] = value;
fill = (fill + 1) % MAX

Thread1:
fill = (fill + 1) % MAX

A solution: adding mutual exclusion

- What we've forgotten here is mutual exclusion
 - The filling of a buffer and incrementing of the index into the buffer is a critical section, and thus must be guarded carefully.
 - Any idea?

```

sem_t empty;    sem_t full;    sem_t mutex;
void *producer(void *arg) {
    for (int i = 0; i < loops; i++) {
        sem_wait(&mutex);
        sem_wait(&empty);
        put(i);
        sem_post(&full);
        sem_post(&mutex);
    }
}

void *consumer(void *arg) {
    while (tmp != -1) {
        sem_wait(&mutex);
        sem_wait(&full);
        tmp = get();
        sem_post(&empty);
        sem_post(&mutex);
    }
}

int main(int argc, char *argv[]) {
    sem_init(&empty, 0, MAX); // MAX buffers are empty to begin
    sem_init(&full, 0, 0); // ... and 0 are full
    sem_init(&mutex, 0, 1);
    ...
}

```

Question:

Any problem of this implementation?

```

sem_t empty;    sem_t full;    sem_t mutex;
void *producer(void *arg) {
    for (int i = 0; i < loops; i++) {
        sem_wait(&mutex);
        sem_wait(&empty);
        put(i);
        sem_post(&full);
        sem_post(&mutex);
    }
}
void *consumer(void *arg) {
    while (tmp != -1) {
        sem_wait(&mutex);
        sem_wait(&full);
        tmp = get();
        sem_post(&empty);
        sem_post(&mutex);
    }
}
int main(int argc, char *argv[]) {
    sem_init(&empty, 0, MAX); // MAX buffers are
    sem_init(&full, 0, 0); // ... and 0 are full
    sem_init(&mutex, 0, 1);
    ...
}

```

Dead lock:

1. A consumer comes first, acquires the **mutex**, call wait(&full), since the buffer is empty, the consumer will be blocked waiting for a producer to post(&full);
2. A producer comes later, try to acquire the **mutex**, but not successful, will be blocked and wait for a consumer to post(&mutex);
3. Ending up with waiting for each other

```

sem_t empty;    sem_t full;    sem_t mutex;
void *producer(void *arg) {
    for (int i = 0; i < loops; i++) {
        sem_wait(&mutex);
        sem_wait(&empty);
        put(i);
        sem_post(&full);
        sem_post(&mutex);
    }
}

void *consumer(void *arg) {
    while (tmp != -1) {
        sem_wait(&mutex);
        sem_wait(&full);
        tmp = get();
        sem_post(&empty);
        sem_post(&mutex);
    }
}

int main(int argc, char *argv[]) {
    sem_init(&empty, 0, MAX); // MAX buffers are empty to begin
    sem_init(&full, 0, 0); // ... and 0 are full
    sem_init(&mutex, 0, 1);
    ...
}

```

How to fix this?

Hint:

Change the
scope of the lock

```

sem_t empty;    sem_t full;    sem_t mutex;
void *producer(void *arg) {
    for (int i = 0; i < loops; i++) {
        sem_wait(&empty);
        sem_wait(&mutex);
        put(i);
        sem_post(&mutex);
        sem_post(&full);
    }
}

void *consumer(void *arg) {
    while (tmp != -1) {
        sem_wait(&full);
        sem_wait(&mutex);
        tmp = get();
        sem_post(&mutex);
        sem_post(&empty);
    }
}

int main(int argc, char *argv[]) {
    sem_init(&empty, 0, MAX); // MAX buffers are empty to begin
    sem_init(&full, 0, 0); // ... and 0 are full
    sem_init(&mutex, 0, 1);
    ...
}

```

How to fix this?

Answer:

Change the
order of the lock

Protect only the
critical section.

Outline

- Why semaphores?
- What is a semaphore?
- How to use semaphores?
 - Binary semaphores (similar to locks)
 - Semaphore for ordering (similar to condition variable)
 - Producer and Consumer using semaphore
 - Implementing semaphore with lock and condition variable
- Java methods

Simulate Semaphores with locks and condition variables

- Example of using locks and condition variable to simulate semaphore, called “Zemaphore”

```
typedef struct __Zem_t {  
    int value;  
    pthread_cond_t cond;  
    pthread_mutex_t lock;  
} Zem_t;
```



```
// only one thread can call this
void Zem_init(Zem_t *s, int value) {
    s->value = value;
    Cond_init(&s->cond);
    Mutex_init(&s->lock);
}
```

```
void Zem_wait(Zem_t *s) {
    Mutex_lock(&s->lock);
    while (s->value <= 0)
        Cond_wait(&s->cond, &s->lock);
    s->value--;
    Mutex_unlock(&s->lock);
}
```

```
void Zem_post(Zem_t *s) {
    Mutex_lock(&s->lock);
    s->value++;
    Cond_signal(&s->cond);
    Mutex_unlock(&s->lock);
}
```

Outline

- Why semaphores?
- What is a semaphore?
- How to use semaphores?
 - Binary semaphores (similar to locks)
 - Semaphore for ordering (similar to condition variable)
 - Producer and Consumer using semaphore
 - Implementing semaphore with lock and condition variable
- Java methods

Java methods

```
Semaphore sem = new Semaphore(count) ;  
//or  
sem = new Semaphore(count, true/false) ;  
// true: strong - false: weak
```

```
void acquire() ; // P
```

```
void release() ; // V
```

```
int availablePermits() ; //returns count
```

```
/* more useful functions in:
```

```
https://docs.oracle.com/javase/10/docs/api/java/util/concurrent/Semaphore.html
```

```
*/
```

Strong/Weak Semaphore

- A **queue** is used to hold threads **waiting** on the semaphore
 - In what order are threads removed from the queue?
- **Strong Semaphores** use **FIFO**
- **Weak Semaphores** not deterministic

Acknowledgement

- Chapter 31
 - Operating Systems: Three Easy Pieces
- Java documentation
 - <https://docs.oracle.com/javase/10/docs/api/java/util/concurrent/Semaphore.html>

Further reading (recommended, not required)

- One great (and free reference) is Allen Downey's book on concurrency and programming with semaphores:
- **"The Little Book of Semaphores"**
- This book has lots of puzzles you can work on to improve your understanding of both semaphores in specific and concurrency in general.

Questions?